

MODULE 3: HOUSING PRICES ASSIGNMENT

Claire Markey, Julia Granito, Manny Hurtado, and Steve Desilets

MSDS 422: Practical Machine Learning

April 16th, 2023

Introduction

Developing a well-tuned predictive model (i.e, Ridge, Lasso, ElasticNet) can provide accurate, precise estimates of home prices, which may be helpful to real-estate industry stakeholders. Conceptually, a Lasso or ElasticNet approach may perform better if there are many good predictors due to cost functions. All of these methods offer advantages such as regularization, reducing collinearities, and improving performance on unseen data. Conversely, there are negative consequences such as not being the best options for variable selection due to the algorithms arbitrarily selecting which variable(s) to allocate to in the presence of multicollinearity. To that end, predictive models (i.e., Ridge, Lasso, ElasticNet) were evaluated to identify predictors of home prices and to determine which approach may be best suited to predict home prices.

Method

Data about home sale prices in Ames, Iowa, were downloaded from Kaggle and analyzed using Jupyter Notebooks (Montoya, 2016). An exploratory data analysis explored feature characteristics (associations, distributions, missingness, and outliers), as well as bivariate (Pearson correlations, t-tests, ANOVAs, simple linear regressions) and multivariate analyses (OLS, Ridge, Lasso, and ElasticNet regression). Multivariate analyses were conducted using a k -fold cross-validation design (James et al. 2021) and Kaggle training/validation sets. Underlying model assumptions and metrics were examined in analyses.

Results and Insights

First, descriptive statistics about the outcome of interest (Sale Price) were examined by constructing histograms and boxplots, then examined by calculating summary statistics. Sale price was skewed right, ranged from \$34,900 to \$755,000 ($M = \$180,921$), and included outliers. A log transformation of the home sale price variable improved the normality of the outcome variable ahead of subsequent analyses. After addressing missing data and outliers, categorical variables were separated into dichotomous indicator, ordinal indicator, and non-indicator variables in order to be encoded. An alternative, larger dataset was created with these new variables. We used this dataset later on in the analysis that explored an ElasticNet regression model.

To determine which features should be included in OLS models, we examined the associations between SalePrice and features. Bivariate analyses suggested the following variables had strong relationships with Sale Price: central air ($t(1458) = 17.26, p < .001$), exterior quality ($F(3, 1456) = 443.33, p < .001$), and total square footage (TotalSF, a linear combination of GrLivArea and TotalBsmtSF; $r = .82, p < .001$). Simple linear regressions were conducted between Sale Price and its key predictors. A simple linear regression, piecewise regression model, and parsimonious five-factor multiple regression model exhibited R^2 values of .76, .82, and .84, respectively. Then a second-degree polynomial regression on principal components and a regression using L2 regularization on principal components were considered, which returned R^2 values of .91.

To further this analysis, regularization techniques were explored (Ridge, Lasso, ElasticNet) on the full dataset; these analyses included all encoded variables. We observe empirically that non-zero levels of alpha resulted in more generalizable models.

Hyperparameter tuning using different methods, including a five-fold cross-validation design, were then conducted to identify the alpha that minimizes the error on the testing data. First, using R^2 as the performance metric, we plotted training and testing data for alphas at even increments from 0 to .02. By

inspection, there was an inflection point for the Lasso regularization technique. We then assessed that the optimal value for alpha was about .004, with an R^2 of .90.

Next, we explored a Lasso regression model to model the log sale prices. To do so, we standardized each of the numeric predictor variables and fit a Lasso regression model to the training data using the optimal tuning parameter. The R squared of this lasso model when applied to a validation set was 0.863. After applying the model to the Kaggle test dataset, the resulting root mean squared error was 0.16194.

For Ridge regression analyses, the optimal alpha was approximately .015 for the initial data; the alpha increased to .81 and the R^2 was .89 after including categorical variables. We found that our ElasticNet model performed the best out of the three regularization methods when all the encoded variables (full data set) were included, with an R^2 for the final model of .94 and a Kaggle RMSE score of .1356. Regarding ElasticNet, the best alpha values were .015 and .009 for the limited and full data sets, respectively. This indicates that a combination of L1 and L2 regularization is optimal for our dataset. The implications of an alpha close to zero for ElasticNet regression means that the importance of the predictors is likely to be distributed among the features. This results in the coefficients of important variables to be close, despite any high correlations. If the alpha value were closer to 1, therefore suggesting higher bias, then the coefficient for one of the correlated variables would be high, while the others would have shrunk to zero. In comparing the coefficients, we observe that the Lasso model has more non-zero coefficients with lower magnitude, while the ElasticNet coefficients are larger but the matrix is somewhat more sparse. By incorporating both L1 and L2 regularization, ElasticNet optimizes the cost function to produce the best fit.

Models such as ElasticNet are not ideal for feature selection as the penalty term may cause the model to skew the weights of certain predictors, making some good predictors appear unimportant. However, we do consider the variables in our model in the analysis. Regarding the Elastic Net model that achieved the lowest Kaggle RMSE score (0.1356), the largest coefficients included square footage variables (Total Square Ft), Overall Quality and Condition variables (overall material, finish, and condition of a home), and the Year Built variable. These had positive coefficients indicating Sale Price increased as the predictors increased. The coefficient of one of the encoded variables, indicating whether the home is in the Crawford neighborhood, was large, indicating that living in this neighborhood was associated with higher Sale Prices. Some variables have negative coefficients, like the engineered feature, Year Since Last remodel, and like some of the of the encoded variables related to the zoning classification of properties, specifically for commercial and medium-density residential properties. This indicates that a greater number of years since a remodel or selling a commercial or medium-density residential property can negatively impact the sale price of a home.

Exploratory data analysis provided insights about factors that impact the sale price of homes in Ames. This work successfully explored methods to accurately predict Sale Price using regression techniques (simple linear, MLR, Piecewise, Polynomial, Ridge, Lasso, and Elastic Net). Our highest performing model, built using Elastic Net regression, has practical implications for Ames real estate market stakeholders. Given the variability of the housing market over time, along with the fact that the housing market in Ames may not be representative of other cities or regions, extrapolating the results beyond the timeframe and city for the data used to train the model will likely lead to inaccurate predictions.

References

- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2021. *An Introduction to Statistical Learning*. Springer. <https://www.statlearning.com/>
- Montoya, Anna. 2016. "House Prices - Advanced Regression Techniques." *Kaggle*. <https://www.kaggle.com/c/house-prices-advanced-regression-techniques>

Appendix - Python Code and Outputs

Data Preparation

First, we will set up this notebook so that it will display multiple outputs for each cell if needed.

```
In [1]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

Second, we will import the data. We will view the first five rows of data and the shape of the dataframe to confirm that the data imported correctly.

```
In [2]: import pandas as pd
housing_training_data = pd.read_csv('train.csv')

# show first five rows of the data
housing_training_data.head()
# show number of columns and rows
housing_training_data.shape
```

```
Out[2]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub

5 rows × 81 columns

```
Out[2]: (1460, 81)
```

Distribution of the Dependent Variable


We can begin examining the distribution of this dataset's dependent variable, sale price, by generating summary statistics for this variable.

```
In [3]: housing_training_data['SalePrice'].describe()
```


```
Out[3]: count      1460.000000  
mean      180921.195890  
std       79442.502883  
min       34900.000000  
25%      129975.000000  
50%      163000.000000  
75%      214000.000000  
max       755000.000000  
Name: SalePrice, dtype: float64
```

We can also construct a histogram and a boxplot to visualize the distribution of the sale price variable in this dataframe.

```
In [4]: import seaborn as sns  
import matplotlib.pyplot as plt  
  
histogram = housing_training_data['SalePrice'].hist(edgecolor = 'black', bins = 15, fi
```



```
In [5]: sns.boxplot(x=housing_training_data["SalePrice"])  
  
Out[5]: <AxesSubplot:xlabel='SalePrice'>
```



```
In [6]: import numpy as np
from scipy import stats
from scipy.stats import norm, kurtosis
df = []
raw_data = housing_training_data['SalePrice']
transform_data = np.log(housing_training_data['SalePrice'])
transform_data2, best_lambda = stats.boxcox(housing_training_data['SalePrice'])

print("homeprice kurtosis:", kurtosis(raw_data))
print("log of homeprice kurtosis:", kurtosis(transform_data))
print("boxcox transform of homeprice kurtosis:", kurtosis(transform_data2))

plt.rcParams["figure.figsize"] = [7.50, 3.50]
plt.rcParams["figure.autolayout"] = True

s1 = pd.DataFrame(raw_data)
s2 = pd.DataFrame(np.array(transform_data).tolist(), columns = ['logSalePrice'])
s3 = pd.DataFrame(np.array(transform_data2).tolist(), columns = ['boxcoxSalePrice'])

fig, axes = plt.subplots(1, 3)

s1.hist('SalePrice', ax=axes[0])
s2.hist('logSalePrice', ax=axes[1])
s3.hist('boxcoxSalePrice', ax=axes[2])
plt.show()
```

homeprice kurtosis: 6.509812011089439

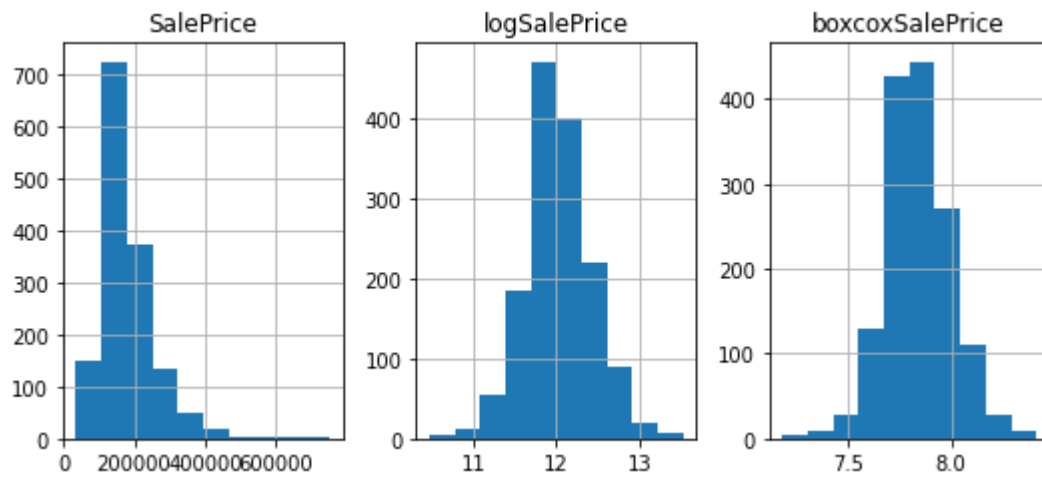
log of homeprice kurtosis: 0.8026555069117713

boxcox transform of homeprice kurtosis: 0.870759906431624

Out[6]: array([<AxesSubplot:title={ 'center': 'SalePrice' }>], dtype=object)

Out[6]: array([<AxesSubplot:title={ 'center': 'logSalePrice' }>], dtype=object)

Out[6]: array([<AxesSubplot:title={ 'center': 'boxcoxSalePrice' }>], dtype=object)



Shapiro Wilk test for normality

```
In [7]: print("Shapiro Wilk test for normality: ", stats.shapiro(raw_data))
        print("Shapiro Wilk test for normality: ", stats.shapiro(transform_data))
        print("Shapiro Wilk test for normality: ", stats.shapiro(transform_data2))
```

```
Shapiro Wilk test for normality: (0.869671642780304, 3.206247534576162e-33)
Shapiro Wilk test for normality: (0.9912067651748657, 1.1490678986092462e-07)
Shapiro Wilk test for normality: (0.9915341138839722, 1.906367685933219e-07)
```

The Shapiro Wilk test for normality (H_0 : normal, H_a : not-normal) suggests a departure from normality for both the raw and transformed data.

Investigation of Missing Data and Outliers

We can take a look at the counts of reported values in each column to determine the number of missing values for each variable in the dataframe.

```
In [8]: # find null counts, percentage of null values, and column type
        null_count = housing_training_data.isnull().sum()
        null_percentage = housing_training_data.isnull().sum() * 100 / len(housing_training_data)
        column_type = housing_training_data.dtypes

        # show null counts, percentage of null values, and column type for columns with more than 1% null
        null_summary = pd.concat([null_count, null_percentage, column_type], axis=1, keys=['Null Count', 'Null Percentage', 'Column Type'])
        null_summary_only_missing = null_summary[null_count != 0].sort_values('Percentage Miss')
        null_summary_only_missing
```


Out[8]:

	Missing Count	Percentage Missing	Column Type
PoolQC	1453	99.520548	object
MiscFeature	1406	96.301370	object
Alley	1369	93.767123	object
Fence	1179	80.753425	object
FireplaceQu	690	47.260274	object
LotFrontage	259	17.739726	float64
GarageType	81	5.547945	object
GarageYrBlt	81	5.547945	float64
GarageFinish	81	5.547945	object
GarageQual	81	5.547945	object
GarageCond	81	5.547945	object
BsmtExposure	38	2.602740	object
BsmtFinType2	38	2.602740	object
BsmtFinType1	37	2.534247	object
BsmtCond	37	2.534247	object
BsmtQual	37	2.534247	object
MasVnrArea	8	0.547945	float64
MasVnrType	8	0.547945	object
Electrical	1	0.068493	object

We will deal with columns that contain missing values. For the purpose of this exploratory data analysis, we will use the percentage of nulls missing, the column type, and the other columns present in the data that may provide information that can be used to fill in the missing values.

We will remove columns with over 50% Null values.

```
In [9]: # PoolQC, MiscFeature, Alley, Fence all have over 50% of missing values, we will remove them
housing_training_data.drop(['Alley', 'PoolQC', 'Fence', 'MiscFeature'], axis=1, inplace=True)

# show new shape
housing_training_data.shape
```

Out[9]: (1460, 77)

We will set Null values in columns that are non-numeric to None.

```
In [10]: # select non-numeric columns that contain more than 1 Null value
columns_None = ['BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'GarageType', 'GarageFinish', 'BsmtExposure', 'BsmtFinType2', 'GarageType', 'GarageFinish']
# set Nulls in non-numeric columns to 'None'
housing_training_data[columns_None] = housing_training_data[columns_None].fillna('None')
```

We determine the best way to handle nulls for each numeric column. We replace nulls in Masonry veneer area with 0, nulls in Lot Frontage with the median, and nulls in Year Garage was built with the average between the year the garage was built and year house was built .

```
In [11]: # change Null values to 0 for Masonry veneer area
housing_training_data['MasVnrArea'].fillna(0, inplace=True)

# show distribution stats for Lot Frontage
housing_training_data['LotFrontage'].describe()
# fill Nulls for Lot Frontage with median value
housing_training_data['LotFrontage'].fillna(housing_training_data['LotFrontage'].median(), inplace=True)

# average years between garage being built and years built
avg_years = round((housing_training_data['GarageYrBlt'] - housing_training_data['YearBuilt'])/2)
# fill Nulls with avg bet year garage was built and year house was built
housing_training_data['GarageYrBlt'].fillna(housing_training_data['YearBuilt']+avg_years, inplace=True)
```

```
Out[11]: count    1201.000000
mean       70.049958
std        24.284752
min        21.000000
25%        59.000000
50%        69.000000
75%        80.000000
max        313.000000
Name: LotFrontage, dtype: float64
```

We can see there are no more missing values in our original dataframe.

```
In [12]: # check that there are no more missing values in the dataframe
null_count = housing_training_data.isnull().sum()
null_count[null_count != 0]
```

```
Out[12]: Series([], dtype: int64)
```

We can also create boxplots for each of the continuous variables in the dataframe to analyze whether outliers exist for each of those variables.

```
In [13]: numerical_vars = ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'MasVnrArea',
                           'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
                           'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
                           'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF',
                           'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'SalePrice']

fig, ax = plt.subplots(17, 2, figsize = (10, 19))

for var, subplot in zip(numerical_vars, ax.flatten()):
    sns.boxplot(x=housing_training_data[var], ax = subplot)

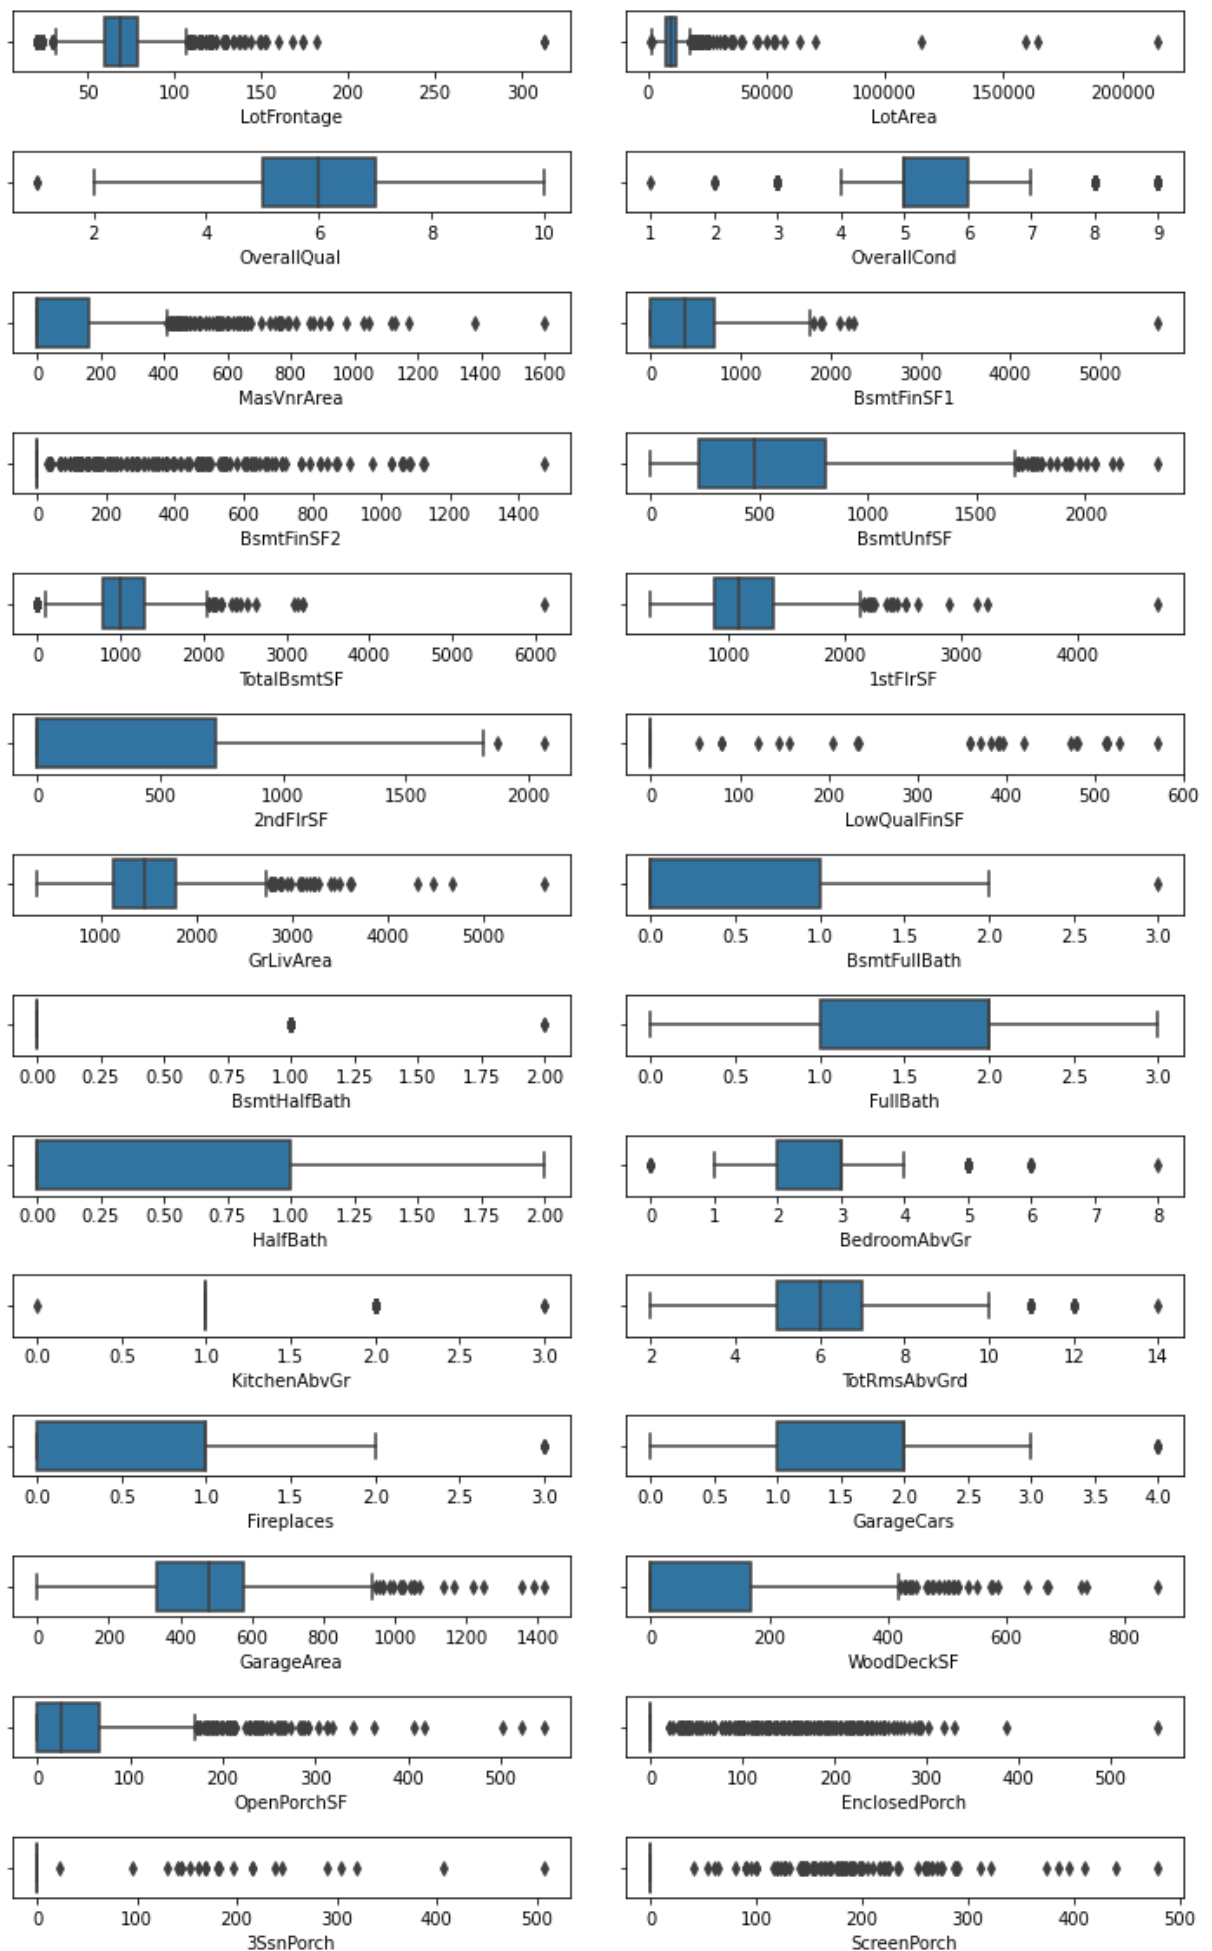
fig.tight_layout()
```

```
Out[13]: <AxesSubplot:xlabel='LotFrontage'>
```

```
Out[13]: <AxesSubplot:xlabel='LotArea'>
```

```
Out[13]: <AxesSubplot:xlabel='OverallQual'>
```

```
Out[13]: <AxesSubplot:xlabel='OverallCond'>
Out[13]: <AxesSubplot:xlabel='MasVnrArea'>
Out[13]: <AxesSubplot:xlabel='BsmtFinSF1'>
Out[13]: <AxesSubplot:xlabel='BsmtFinSF2'>
Out[13]: <AxesSubplot:xlabel='BsmtUnfSF'>
Out[13]: <AxesSubplot:xlabel='TotalBsmtSF'>
Out[13]: <AxesSubplot:xlabel='1stFlrSF'>
Out[13]: <AxesSubplot:xlabel='2ndFlrSF'>
Out[13]: <AxesSubplot:xlabel='LowQualFinSF'>
Out[13]: <AxesSubplot:xlabel='GrLivArea'>
Out[13]: <AxesSubplot:xlabel='BsmtFullBath'>
Out[13]: <AxesSubplot:xlabel='BsmtHalfBath'>
Out[13]: <AxesSubplot:xlabel='FullBath'>
Out[13]: <AxesSubplot:xlabel='HalfBath'>
Out[13]: <AxesSubplot:xlabel='BedroomAbvGr'>
Out[13]: <AxesSubplot:xlabel='KitchenAbvGr'>
Out[13]: <AxesSubplot:xlabel='TotRmsAbvGrd'>
Out[13]: <AxesSubplot:xlabel='Fireplaces'>
Out[13]: <AxesSubplot:xlabel='GarageCars'>
Out[13]: <AxesSubplot:xlabel='GarageArea'>
Out[13]: <AxesSubplot:xlabel='WoodDeckSF'>
Out[13]: <AxesSubplot:xlabel='OpenPorchSF'>
Out[13]: <AxesSubplot:xlabel='EnclosedPorch'>
Out[13]: <AxesSubplot:xlabel='3SsnPorch'>
Out[13]: <AxesSubplot:xlabel='ScreenPorch'>
Out[13]: <AxesSubplot:xlabel='PoolArea'>
Out[13]: <AxesSubplot:xlabel='MiscVal'>
Out[13]: <AxesSubplot:xlabel='SalePrice'>
```




Examination of the Relationship between the Dependent Variable and Potential Predictors

We can use a correlation heatmap to quantify the correlation between the dependent variable, sale price, and the potential continuous predictor variables.

```
In [14]: df_corr_housing_training = housing_training_data[numerical_vars]
          corrmatrix_housing_training = df_corr_housing_training.corr()

          f, ax = plt.subplots(figsize = (20, 20))
          sns.heatmap(corrmatrix_housing_training, vmax = 1, square = True, annot = True,
                      cmap = 'RdYlBu', linewidths = 0.5, fmt=".1f")
```

```
Out[14]: <AxesSubplot:>
```



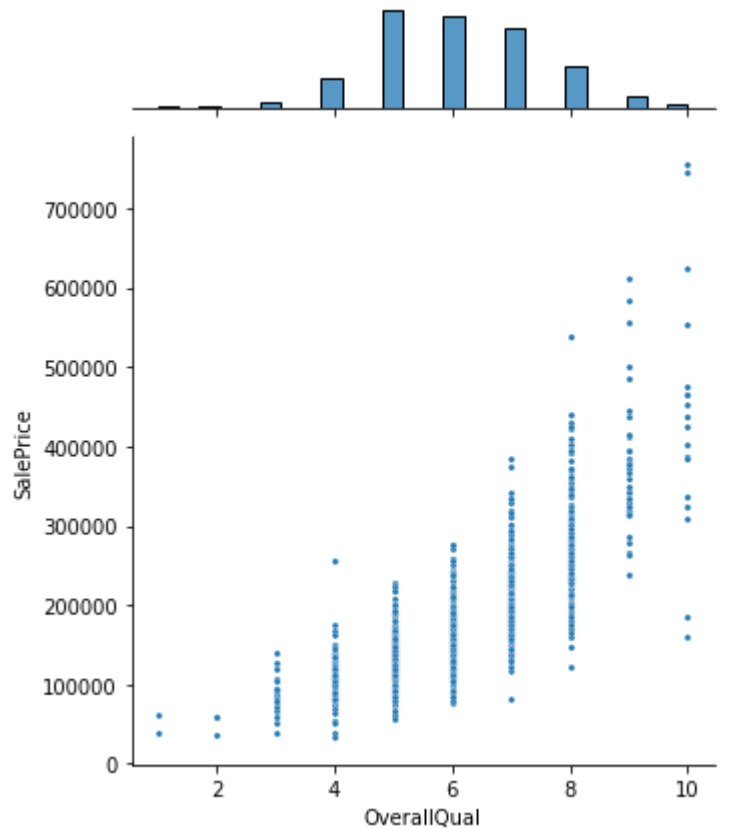
```
In [15]: #Correlation with output variable
cor_target = abs(corrmat_housing_training["SalePrice"])
#Selecting highly correlated features
relevant_features = cor_target[cor_target>0.5]
relevant_features.sort_values(ascending=False)
```

```
Out[15]: SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
Name: SalePrice, dtype: float64
```


We can use jointplots to take a closer look at the relationship between sale price and five of the continuous variables with which sale price has a strong or moderate association: OverallQual, GrLivArea, GarageArea, Fullbath, and TotalBsmntSF.

Below are plots that examine the relationship between variables of interest and sale price


```
In [16]: sns.jointplot(x='OverallQual', y='SalePrice', data = housing_training_data, joint_kws=
```




```
In [17]: sns.jointplot(x='GrLivArea', y='SalePrice', data = housing_training_data, joint_kws=
```




In [18]: `sns.jointplot(x='FullBath', y='SalePrice', data = housing_training_data, joint_kws={"s`



In [19]: `sns.jointplot(x='TotalBsmtSF', y='SalePrice', data = housing_training_data, joint_kws=`



In [20]: `sns.jointplot(x='GarageArea', y='SalePrice', data = housing_training_data, joint_kws={`



To determine which binary categorical variables might serve as the best predictors in a regression model, we can create boxplots and run t-tests to help decipher which binary indicator variables may have the strongest relationship with home sale prices.

```
In [21]: categorical_variables = ['MSZoning', 'Street', 'LotShape', 'LandContour', 'Utilities',  
                                'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'House',  
                                'Exterior1st', 'Exterior2nd', 'MasVnrType', 'ExterQual', 'Exter',  
                                'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtF',  
                                'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'Fire',  
                                'GarageQual', 'GarageCond', 'PavedDrive', 'SaleType', 'SaleCo',  
                                'YrSold', 'MoSold']  
  
category_counts = []  
  
for var in categorical_variables:  
    category_counts.append(len(housing_training_data[var].unique()))  
  
categorical_variable_dictionary = {'Categorical Predictor':categorical_variables,'Number of Categories':category_counts}  
categorical_var_df = pd.DataFrame(categorical_variable_dictionary)  
categorical_var_df  
  
# Identify the Indicator Variables  
indicator_variables_df = categorical_var_df[categorical_var_df['Number of Categories'] > 1]  
indicator_variables_df  
  
# Identify the Non-Indicator Categorical Variables  
non_indicator_categorical_vars_df = categorical_var_df[categorical_var_df['Number of Categories'] == 1]  
non_indicator_categorical_vars_df
```

Out[21]:

	Categorical Predictor	Number of Categories
0	MSZoning	5
1	Street	2
2	LotShape	4
3	LandContour	4
4	Utilities	2
5	LotConfig	5
6	LandSlope	3
7	Neighborhood	25
8	Condition1	9
9	Condition2	8
10	BldgType	5
11	HouseStyle	8
12	RoofStyle	6
13	RoofMatl	8
14	Exterior1st	15
15	Exterior2nd	16
16	MasVnrType	4
17	ExterQual	4
18	ExterCond	5
19	Foundation	6
20	BsmtQual	5
21	BsmtCond	5
22	BsmtExposure	5
23	BsmtFinType1	7
24	BsmtFinType2	7
25	Heating	6
26	HeatingQC	5
27	CentralAir	2
28	Electrical	6
29	KitchenQual	4
30	Functional	7
31	FireplaceQu	6
32	GarageType	7
33	GarageQual	6

	Categorical Predictor	Number of Categories
34	GarageCond	6
35	PavedDrive	3
36	SaleType	9
37	SaleCondition	6
38	YearBuilt	112
39	GarageYrBlt	102
40	YrSold	5
41	MoSold	12

Out[21]:

	Categorical Predictor	Number of Categories
1	Street	2
4	Utilities	2
27	CentralAir	2

Out[21]:

	Categorical Predictor	Number of Categories
0	MSZoning	5
2	LotShape	4
3	LandContour	4
5	LotConfig	5
6	LandSlope	3
7	Neighborhood	25
8	Condition1	9
9	Condition2	8
10	BldgType	5
11	HouseStyle	8
12	RoofStyle	6
13	RoofMatl	8
14	Exterior1st	15
15	Exterior2nd	16
16	MasVnrType	4
17	ExterQual	4
18	ExterCond	5
19	Foundation	6
20	BsmtQual	5
21	BsmtCond	5
22	BsmtExposure	5
23	BsmtFinType1	7
24	BsmtFinType2	7
25	Heating	6
26	HeatingQC	5
28	Electrical	6
29	KitchenQual	4
30	Functional	7
31	FireplaceQu	6
32	GarageType	7
33	GarageQual	6
34	GarageCond	6
35	PavedDrive	3
36	SaleType	9

	Categorical Predictor	Number of Categories
37	SaleCondition	6
38	YearBuilt	112
39	GarageYrBlt	102
40	YrSold	5
41	MoSold	12

```
In [22]: # convert Paved Drive to dichotomous, indicator variable
housing_training_data['PavedDrive'] = np.where(housing_training_data['PavedDrive'] ==
housing_training_data['PavedDrive'].value_counts()

# view indicator variables
indicator_vars = ['Street', 'Utilities', 'CentralAir', 'PavedDrive']

fig, ax = plt.subplots(1, 4, figsize=(15, 5))

for var, subplot in zip(indicator_vars, ax.flatten()):
    sns.boxplot(x = var, y = 'SalePrice', data=housing_training_data, ax=subplot)

fig.tight_layout()
```

```
Out[22]: Y    1340
N     120
```


```
Name: PavedDrive, dtype: int64
```

```
Out[22]: <AxesSubplot:xlabel='Street', ylabel='SalePrice'>
```

```
Out[22]: <AxesSubplot:xlabel='Utilities', ylabel='SalePrice'>
```

```
Out[22]: <AxesSubplot:xlabel='CentralAir', ylabel='SalePrice'>
```

```
Out[22]: <AxesSubplot:xlabel='PavedDrive', ylabel='SalePrice'>
```



```
In [23]: # Run T-Tests To Determine Which Indicator Variables Might Have the Strongest Association
from scipy.stats import ttest_ind

Street_t_test = ttest_ind(housing_training_data['SalePrice'][housing_training_data['Street'] == 'Grvl'],
housing_training_data['SalePrice'][housing_training_data['Street'] == 'Pave'],
equal_var=False)

Utilities_t_test = ttest_ind(housing_training_data['SalePrice'][housing_training_data['Utilities'] == 'AllPub'],
housing_training_data['SalePrice'][housing_training_data['Utilities'] == 'NoSeWa'],
equal_var=False)
```

```

Central_Air_t_test = ttest_ind(housing_training_data['SalePrice'][housing_training_data['CentralAir'] == 'No'],
                               housing_training_data['SalePrice'][housing_training_data['CentralAir'] == 'Yes'],
                               equal_var=False)

Paved_Drive_t_test = ttest_ind(housing_training_data['SalePrice'][housing_training_data['PavedDrive'] == 'No'],
                               housing_training_data['SalePrice'][housing_training_data['PavedDrive'] == 'Yes'],
                               equal_var=False)

Indicator_Variable_t_test_statistics = [Street_t_test[0], Utilities_t_test[0], Central_Air_t_test[0]]
Indicator_Variable_t_test_p_values = [Street_t_test[1], Utilities_t_test[1], Central_Air_t_test[1]]

indicator_var_t_tests = {'Indicator Variable': indicator_vars, 'T-Test Statistic': Indicator_Variable_t_test_statistics,
                        'P-Values': Indicator_Variable_t_test_p_values}
Indicator_var_t_test_df = pd.DataFrame(indicator_var_t_tests)
Indicator_var_t_test_df.style.background_gradient(cmap = 'Greens')

```

```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/numpy/core/fromnumeric.py:3724: RuntimeWarning: Degrees of freedom <= 0 for slice
  **kwargs)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/numpy/core/_methods.py:254: RuntimeWarning: invalid value encountered in double_scalars
  ret = ret.dtype.type(ret / rcount)

```

Out[23]:

	Indicator Variable	T-Test Statistic	P-Values
0	Street	1.900788	0.115048
1	Utilities	nan	nan
2	CentralAir	17.267773	0.000000
3	PavedDrive	15.093535	0.000000

We can see that Street, Central Air, and Paved Drive all have statistically significant t-test statistics. This tells us that there does appear to be statistically significant differences in Sale Prices between the categories of these variables. Let's dummy encode each one so that we can use them in our regression analysis.

In []:

```

In [25]: # dummy encode the Street, Central Air, Paved Drive indicator variables, we will exclude
housing_training_data = pd.get_dummies(housing_training_data, columns=['Street', 'CentralAir', 'PavedDrive'])

```

To determine which categorical variables might be most useful for inclusion in a regression model (in the form of a dichotomous variable), we can create boxplots and run analyses of variance (ANOVA) to determine which non-binary categorical variables may have the strongest relationship with home sale prices.

We can create boxplots to visually display the distribution of sale prices disaggregated by the categories associated with each of the non-indicator categorical variables as well.

```
In [26]: # redefine categorical_vars_df after transforming Paved Drive variable
non_indicator_categorical_vars_df = non_indicator_categorical_vars_df[non_indicator_cate

fig, ax = plt.subplots(33, 1, figsize=(15, 200))


for var, subplot in zip(non_indicator_categorical_vars_df['Categorical Predictor'], ax,
                        sns.boxplot(x = var, y = 'SalePrice', data=housing_training_data, ax=subplot)

fig.tight_layout()
```

```
Out[26]: <AxesSubplot:xlabel='MSZoning', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='LotShape', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='LandContour', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='LotConfig', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='LandSlope', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Neighborhood', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Condition1', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Condition2', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BldgType', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='HouseStyle', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='RoofStyle', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='RoofMatl', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Exterior1st', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Exterior2nd', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='MasVnrType', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='ExterQual', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='ExterCond', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Foundation', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BsmtQual', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BsmtCond', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BsmtExposure', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BsmtFinType1', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='BsmtFinType2', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Heating', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='HeatingQC', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='Electrical', ylabel='SalePrice'>
Out[26]: <AxesSubplot:xlabel='KitchenQual', ylabel='SalePrice'>
```



```
Out[26]: <AxesSubplot:xlabel='Functional', ylabel='SalePrice'>  
Out[26]: <AxesSubplot:xlabel='FireplaceQu', ylabel='SalePrice'>  
Out[26]: <AxesSubplot:xlabel='GarageType', ylabel='SalePrice'>  
Out[26]: <AxesSubplot:xlabel='GarageQual', ylabel='SalePrice'>  
Out[26]: <AxesSubplot:xlabel='GarageCond', ylabel='SalePrice'>  
Out[26]: <AxesSubplot:xlabel='SaleType', ylabel='SalePrice'>
```



Visual inspection of the boxplots above suggests that the variables for exterior quality, basement quality, fireplace quality, kitchen quality, garage quality and condition, and Heating quality may provide the most promise in our search for helpful categorical predictors that may be transformed into dichotomous variables. Conducting ANOVAs can shed more light on the relationship between these seven variables and home sale prices.

```
In [27]: ANOVA_variables = ['ExterQual', 'BsmtQual', 'FireplaceQu', 'KitchenQual', 'GarageQual',

from scipy.stats import f_oneway

# ExterQual ANOVA
ExterQual_Gd = housing_training_data['SalePrice'][housing_training_data['ExterQual'] == 'Gd']
ExterQual_TA = housing_training_data['SalePrice'][housing_training_data['ExterQual'] == 'TA']
ExterQual_Ex = housing_training_data['SalePrice'][housing_training_data['ExterQual'] == 'Ex']
ExterQual_Fa = housing_training_data['SalePrice'][housing_training_data['ExterQual'] == 'Fa']

ANOVA_ExterQual = f_oneway(ExterQual_Gd, ExterQual_TA, ExterQual_Ex, ExterQual_Fa)

# BsmtQual ANOVA
BsmtQual_Gd = housing_training_data['SalePrice'][housing_training_data['BsmtQual'] == 'Gd']
BsmtQual_TA = housing_training_data['SalePrice'][housing_training_data['BsmtQual'] == 'TA']
BsmtQual_Ex = housing_training_data['SalePrice'][housing_training_data['BsmtQual'] == 'Ex']
BsmtQual_None = housing_training_data['SalePrice'][housing_training_data['BsmtQual'] == 'None']
BsmtQual_Fa = housing_training_data['SalePrice'][housing_training_data['BsmtQual'] == 'Fa']

ANOVA_BsmtQual = f_oneway(BsmtQual_Gd, BsmtQual_TA, BsmtQual_Ex, BsmtQual_None, BsmtQual_Fa)

# FireplaceQu ANOVA
FireplaceQu_None = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'None']
```

```

FireplaceQu_TA = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'TA']
FireplaceQu_Gd = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'Gd']
FireplaceQu_Fa = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'Fa']
FireplaceQu_Ex = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'Ex']
FireplaceQu_Po = housing_training_data['SalePrice'][housing_training_data['FireplaceQu'] == 'Po']

ANOVA_FireplaceQu = f_oneway(FireplaceQu_None, FireplaceQu_TA, FireplaceQu_Gd, FireplaceQu_Fa,
                             FireplaceQu_Ex, FireplaceQu_Po)

# KitchenQual ANOVA
KitchenQual_Gd = housing_training_data['SalePrice'][housing_training_data['KitchenQual'] == 'Gd']
KitchenQual_TA = housing_training_data['SalePrice'][housing_training_data['KitchenQual'] == 'TA']
KitchenQual_Ex = housing_training_data['SalePrice'][housing_training_data['KitchenQual'] == 'Ex']
KitchenQual_Fa = housing_training_data['SalePrice'][housing_training_data['KitchenQual'] == 'Fa']

ANOVA_KitchenQual = f_oneway(KitchenQual_Gd, KitchenQual_TA, KitchenQual_Ex, KitchenQual_Fa,
                             KitchenQual_Po)

# GarageQual ANOVA
GarageQu_None = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'None']
GarageQu_TA = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'TA']
GarageQu_Gd = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'Gd']
GarageQu_Fa = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'Fa']
GarageQu_Ex = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'Ex']
GarageQu_Po = housing_training_data['SalePrice'][housing_training_data['GarageQual'] == 'Po']

ANOVA_GarageQu = f_oneway(GarageQu_None, GarageQu_TA, GarageQu_Gd, GarageQu_Fa,
                          GarageQu_Ex, GarageQu_Po)

# GarageCond ANOVA
GarageCond_None = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'None']
GarageCond_TA = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'TA']
GarageCond_Gd = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'Gd']
GarageCond_Fa = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'Fa']
GarageCond_Ex = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'Ex']
GarageCond_Po = housing_training_data['SalePrice'][housing_training_data['GarageCond'] == 'Po']

ANOVA_GarageCond = f_oneway(GarageCond_None, GarageCond_TA, GarageCond_Gd, GarageCond_Fa,
                            GarageCond_Ex, GarageCond_Po)

# HeatingQC ANOVA
HeatingQC_TA = housing_training_data['SalePrice'][housing_training_data['HeatingQC'] == 'TA']
HeatingQC_Gd = housing_training_data['SalePrice'][housing_training_data['HeatingQC'] == 'Gd']
HeatingQC_Fa = housing_training_data['SalePrice'][housing_training_data['HeatingQC'] == 'Fa']
HeatingQC_Ex = housing_training_data['SalePrice'][housing_training_data['HeatingQC'] == 'Ex']
HeatingQC_Po = housing_training_data['SalePrice'][housing_training_data['HeatingQC'] == 'Po']

ANOVA_HeatingQC = f_oneway(HeatingQC_TA, HeatingQC_Gd, HeatingQC_Fa,
                           HeatingQC_Ex, HeatingQC_Po)

# Compile Outputs
ANOVA_statistics = [ANOVA_ExterQual[0], ANOVA_BsmtQual[0], ANOVA_FireplaceQu[0], ANOVA_KitchenQual[0],
                   ANOVA_GarageQu[0], ANOVA_GarageCond[0], ]

ANOVA_p_values = [ANOVA_ExterQual[1], ANOVA_BsmtQual[1], ANOVA_FireplaceQu[1], ANOVA_KitchenQual[1],
                 ANOVA_GarageQu[1], ANOVA_GarageCond[1]]

ANOVA_outputs = {'Categorical Variable': ANOVA_variables, 'Test Statistic': ANOVA_statistics,
                 'P-Values': ANOVA_p_values}

```

```
ANOVA_df = pd.DataFrame(ANOVA_outputs)
ANOVA_df.style.background_gradient(cmap = 'Greens')
```

Out[27]:

	Categorical Variable	Test Statistic	P-Values
0	ExterQual	443.334831	0.000000
1	BsmtQual	316.148635	0.000000
2	FireplaceQu	121.075121	0.000000
3	KitchenQual	407.806352	0.000000
4	GarageQual	88.394462	0.000000
5	GarageCond	25.776093	0.000000
6	HeatingQC	25.750153	0.000000

We will use the Tukey-Cramer Multiple Comparison Test to confirm whether there are statistically significant differences in means when considering pairwise comparisons of categorical variable values.

```
In [28]: # Lets view the tukeyhsd for each variable
from statsmodels.stats.multicomp import pairwise_tukeyhsd

for i in ANOVA_variables:
    tukey_cramer_result = pairwise_tukeyhsd(endog=housing_training_data['SalePrice'],
    print(tukey_cramer_result)
```

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-279375.7473	0.001	-323897.1579	-234854.3366	True
Ex	Gd	-135727.4513	0.001	-157297.3415	-114157.5611	True
Ex	TA	-223019.6481	0.001	-244104.9402	-201934.356	True
Fa	Gd	143648.296	0.001	103567.1096	183729.4823	True
Fa	TA	56356.0992	0.0016	16533.607	96178.5913	True
Gd	TA	-87292.1968	0.001	-95594.9096	-78989.484	True

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-211349.0128	0.001	-241850.1706	-180847.8549	True
Ex	Gd	-124352.5624	0.001	-140151.0688	-108554.0559	True
Ex	None	-221388.1494	0.001	-251243.057	-191533.2419	True
Ex	TA	-186281.2231	0.001	-202017.8288	-170544.6174	True
Fa	Gd	86996.4504	0.001	59383.7383	114609.1625	True
Fa	None	-10039.1367	0.9	-47511.5875	27433.3141	False
Fa	TA	25067.7896	0.095	-2509.553	52645.1322	False
Gd	None	-97035.5871	0.001	-123932.7383	-70138.4358	True
Gd	TA	-61928.6608	0.001	-70860.7377	-52996.5838	True
None	TA	35106.9263	0.0034	8246.0868	61967.7658	True

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-170414.0152	0.001	-221601.4617	-119226.5687	True
Ex	Gd	-111361.0842	0.001	-151519.9919	-71202.1765	True
Ex	None	-196381.0174	0.001	-236000.3876	-156761.6472	True
Ex	Po	-207948.35	0.001	-265717.2832	-150179.4168	True
Ex	TA	-131989.0112	0.001	-172402.4484	-91575.574	True
Fa	Gd	59052.9309	0.001	24425.9294	93679.9324	True
Fa	None	-25967.0022	0.2479	-59966.7943	8032.7899	False
Fa	Po	-37534.3348	0.3543	-91604.1139	16535.4442	False
Fa	TA	38425.004	0.0213	3503.1301	73346.8778	True
Gd	None	-85019.9332	0.001	-97208.8236	-72831.0427	True
Gd	Po	-96587.2658	0.001	-140360.8217	-52813.7099	True
Gd	TA	-20627.927	0.001	-35192.2827	-6063.5712	True
None	Po	-11567.3326	0.9	-54846.4364	31711.7712	False
None	TA	64392.0062	0.001	51389.0416	77394.9708	True
Po	TA	75959.3388	0.001	31952.1549	119966.5227	True

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-222989.4649	0.001	-251454.2769	-194524.6528	True
Ex	Gd	-116438.6461	0.001	-132752.1208	-100125.1714	True
Ex	TA	-188592.1584	0.001	-204662.7844	-172521.5325	True
Fa	Gd	106550.8188	0.001	81616.8092	131484.8283	True
Fa	TA	34397.3064	0.0021	9621.5039	59173.109	True
Gd	TA	-72153.5123	0.001	-80503.6215	-63803.4031	True

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
--------	--------	----------	-------	-------	-------	--------

Ex	Fa	-117426.6458	0.101	-246946.454	12093.1623	False
Ex	Gd	-25139.2857	0.9	-163601.7848	113323.2134	False
Ex	None	-137682.716	0.0264	-265641.1324	-9724.2997	True
Ex	Po	-140833.3333	0.2107	-318533.0441	36866.3774	False
Ex	TA	-53510.164	0.8078	-179306.5197	72286.1917	False
Fa	Gd	92287.3601	0.001	26180.9523	158393.768	True
Fa	None	-20256.0702	0.6669	-59898.8479	19386.7075	False
Fa	Po	-23406.6875	0.9	-152926.4956	106113.1206	False
Fa	TA	63916.4818	0.001	31933.4144	95899.5493	True
Gd	None	-112543.4303	0.001	-175535.7535	-49551.1071	True
Gd	Po	-115694.0476	0.1625	-254156.5467	22768.4515	False
Gd	TA	-28370.8783	0.7094	-86846.5105	30104.7539	False
None	Po	-3150.6173	0.9	-131109.0337	124807.7991	False
None	TA	84172.5521	0.001	59254.8429	109090.2612	True
Po	TA	87323.1693	0.3544	-38473.1863	213119.525	False

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-9345.9714	0.9	-167580.7655	148888.8227	False
Ex	Gd	55930.0	0.9	-114211.491	226071.491	False
Ex	None	-20682.716	0.9	-176469.8843	135104.4522	False
Ex	Po	-15500.0	0.9	-190004.8015	159004.8015	False
Ex	TA	63885.7353	0.8252	-90129.0525	217900.5231	False
Fa	Gd	65275.9714	0.1987	-16067.2196	146619.1625	False
Fa	None	-11336.7446	0.9	-55362.0733	32688.5841	False
Fa	Po	-6154.0286	0.9	-96267.9206	83959.8634	False
Fa	TA	73231.7067	0.001	35960.5027	110502.9107	True
Gd	None	-76612.716	0.0493	-153085.6293	-139.8028	True
Gd	Po	-71430.0	0.4306	-181113.1364	38253.1364	False
Gd	TA	7955.7353	0.9	-64838.6294	80750.1	False
None	Po	5182.716	0.9	-80560.5359	90925.968	False
None	TA	84568.4513	0.001	59657.924	109478.9787	True
Po	TA	79385.7353	0.0672	-3093.4641	161864.9347	False

Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Ex	Fa	-90994.9394	0.001	-119740.4887	-62249.39	True
Ex	Gd	-58055.5578	0.001	-72506.7051	-43604.4104	True
Ex	Po	-127914.4291	0.3799	-322924.5088	67095.6505	False
Ex	TA	-72551.553	0.001	-84383.0712	-60720.0347	True
Fa	Gd	32939.3816	0.0271	2400.2524	63478.5107	True
Fa	Po	-36919.4898	0.9	-233776.6307	159937.6511	False
Fa	TA	18443.3864	0.4273	-10946.8769	47833.6496	False
Gd	Po	-69858.8714	0.8508	-265141.3923	125423.6495	False
Gd	TA	-14495.9952	0.0864	-30190.4728	1198.4824	False
Po	TA	55362.8762	0.9	-139743.2799	250469.0322	False

Given that the ANOVA and Tukey-Cramer tests the levels of these variables appear to be statistically significant from one another (for the most part). This demonstrates that it makes sense to ordinally encode these with HeatingQC as the exception.

This also reveals that sale prices for rows with 'None' is not statistically different from sale prices of rows with 'FA'. For each variable, so we will encode 'None' with the same value as 'FA'.

```
In [29]: # for heatingQC excellent is the only value that is statistically significant from the
housing_training_data['HeatingEx'] = np.where(housing_training_data['HeatingQC'] == 'Ex',
housing_training_data.drop(columns=['HeatingQC'], inplace=True)
```

Encode important categorical variables

```
In [30]: from sklearn.preprocessing import LabelEncoder
# Let's ordinally encode the other variables we explored above (aside from heatingQC w
important_categorical = ['ExterQual', 'BsmtQual', 'FireplaceQu', 'KitchenQual', 'Garag

ordinal_mapping = {
    'Ex': 4,
    'Gd': 3,
    'TA': 2,
    'Fa': 1,
    'None': 1,
    'Po': 0,
}

# process columns, replace to categorical features with ordinal ranking
for i in important_categorical:
    housing_training_data[i] = housing_training_data[i].replace(ordinal_mapping)

# shape
print('Shape all_data: {}'.format(housing_training_data.shape))
```

Shape all_data: (1460, 77)

Feature Creation

New features may enable us to create more accurate prediction models for home sale prices. Accordingly, we will create a feature to reflect the number of years since a home has been remodeled.


```
In [31]: # create new variable, years since the house has been remodeled from selling date (use
housing_training_data['YrSinceRemod'] = housing_training_data['YrSold'] - housing_trai
housing_training_data['YrSinceRemod'].describe()
# create boxplot of YrSinceRemod
sns.boxplot(x = 'YrSinceRemod', data=housing_training_data)

#Pearson correlation coefficient and p value for sale price and GarageArea (Size of ga
res6 = stats.pearsonr(housing_training_data.YrSinceRemod, housing_training_data.SalePr
print("Pearson correlation coefficient and p value for sale price and Years since Hous
res6
```

```
Out[31]: count    1460.000000
         mean      22.950000
         std       20.640653
         min       -1.000000
         25%        4.000000
         50%       14.000000
         75%       41.000000
         max       60.000000
         Name: YrSinceRemod, dtype: float64
Out[31]: <AxesSubplot:xlabel='YrSinceRemod'>
```


Pearson correlation coefficient and p value for sale price and Years since House was remodeled/built:

```
Out[31]: (-0.509078738015629, 4.3748554463775595e-97)
```



We can create a scatterplot to visualize the relationship between years since remodel and sale price.


```
In [32]: sns.jointplot(x='YrSinceRemod', y='SalePrice', data = housing_training_data, joint_kws
```

We will also create a feature to reflect the number of total square feet in a home.

```
In [33]: # create new variable TotalSF
housing_training_data['TotalSF'] = housing_training_data['TotalBsmntSF'] + housing_traini
housing_training_data['TotalSF'].describe()
# create boxplot of TotalSF
sns.boxplot(x = 'TotalSF', data=housing_training_data)
```

```
Out[33]: count      1460.000000
mean        2572.893151
std          823.598492
min           334.000000
25%          2014.000000
50%          2479.000000
75%          3008.500000
max         11752.000000
Name: TotalSF, dtype: float64
<AxesSubplot:xlabel='TotalSF'>
```




```
In [34]: # drop large outlier from the dataframe
housing_training_data.drop(housing_training_data[housing_training_data['TotalSF'] > 6000])
# visualize distribution without extreme outliers
sns.histplot(data=housing_training_data, x="TotalSF")

#Pearson correlation coefficient and p value for sale price and TotalSF):
res7 = stats.pearsonr(housing_training_data.TotalSF, housing_training_data.SalePrice)
print("Pearson correlation coefficient and p value for sale price and TotalSF (Total square feet - includes basement):")
res7
```

Out[34]: <AxesSubplot:xlabel='TotalSF', ylabel='Count'>


Pearson correlation coefficient and p value for sale price and TotalSF (Total square feet - includes basement):

Out[34]: (0.81999629129728, 0.0)



We can create a scatterplot to examine the relationship between total square feet and sale price.

```
In [35]: sns.jointplot(x='TotalSF', y='SalePrice', data = housing_training_data, joint_kws={"s"
```



Dummy Encoding

Let's dummy encode the remaining categorical variables and create a new dataframe with these encoded columns and the original numeric columns.

```
In [36]: # create new df with current numeric columns
housing_training_numeric_df = pd.DataFrame(housing_training_data.select_dtypes(exclude=

# create new df with current categorical columns
housing_training_object_df = pd.DataFrame(housing_training_data.select_dtypes(exclude=
orig_object_cols = housing_training_object_df.columns

# dummy encode the categorical columns
housing_training_data_dummy = pd.get_dummies(housing_training_object_df, columns=orig_

# create new df with original numeric and new dummy encoded columns
housing_training_data_large = pd.concat([housing_training_numeric_df, housing_training
housing_training_data_large
```

Out[36]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1	60	65.0	8450	7	5	2003	2003
1	2	20	80.0	9600	6	8	1976	1976
2	3	60	68.0	11250	7	5	2001	2002
3	4	70	60.0	9550	7	5	1915	1970
4	5	60	84.0	14260	8	5	2000	2000
...
1455	1456	60	62.0	7917	6	5	1999	2000
1456	1457	20	85.0	13175	6	6	1978	1988
1457	1458	70	66.0	9042	7	9	1941	2006
1458	1459	20	68.0	9717	5	6	1950	1996
1459	1460	20	75.0	9937	5	6	1965	1965

1455 rows × 256 columns

In [37]:

```

# Correlation with Sale Price
corrmat_housing_training = housing_training_data_large.corr()
cor_target = abs(corrmat_housing_training["SalePrice"])

# Selecting correlated features
relevant_features = cor_target[cor_target>0.25]

# display features that have correlations with SalePrice over 0.25
with pd.option_context("display.max_rows",300):
    relevant_features.sort_values(ascending=False)

# create new dataframe with only those features
feature_list = relevant_features.index
housing_training_data_large_subset = housing_training_data_large[feature_list]
housing_training_data_large_subset.shape

```

```

Out[37]: SalePrice      1.000000
         TotalSF        0.819996
         OverallQual    0.801196
         GrLivArea      0.718340
         ExterQual       0.695051
         KitchenQual     0.667111
         BsmtQual        0.657389
         GarageCars      0.650951
         TotalBsmtSF     0.643275
         GarageArea      0.638538
         1stFlrSF        0.621678
         FullBath        0.556516
         YearBuilt       0.535781
         TotRmsAbvGrd    0.535425
         FireplaceQu     0.525552
         YrSinceRemod    0.524075
         YearRemodAdd    0.522479
         GarageYrBlt     0.516351
         Foundation_PConc 0.505411
         MasVnrArea      0.476732
         Fireplaces      0.467166
         HeatingEx       0.440949
         BsmtFinType1_GLQ 0.434806
         GarageFinish_Fin 0.422119
         Neighborhood_NridgHt 0.421602
         GarageFinish_Unf 0.417580
         BsmtFinSF1      0.393687
         MasVnrType_None 0.383184
         SaleType_New    0.379298
         SaleCondition_Partial 0.373552
         GarageType_Detchd 0.361482
         MasVnrType_Stone 0.350892
         Foundation_CBlock 0.346991
         GarageType_Attchd 0.340156
         LotFrontage     0.339451
         OpenPorchSF     0.330605
         Exterior2nd_VinylSd 0.328221
         Exterior1st_VinylSd 0.326943
         WoodDeckSF      0.320869
         BsmtExposure_Gd 0.309613
         2ndFlrSF        0.300248
         MSZoning_RM     0.294802
         HalfBath        0.285725
         Neighborhood_NoRidge 0.280726
         GarageQual      0.280230
         LotArea         0.270350
         LotShape_Reg    0.263206
         BsmtExposure_No 0.262934
         SaleType_WD     0.258981
         CentralAir_Y    0.258415
         GarageCond      0.258367
         Name: SalePrice, dtype: float64
Out[37]: (1455, 51)

```

We may use this dataframe later on in the analysis. For the start, we will investigate the relationship between SalePrice and a select number of features, like Total Square Feet.

Prepare Test Data

```
In [38]: # Load test data
housing_testing_data = pd.read_csv('test.csv')
```

Handle Null values, matches how we dealt with Nulls in the training dataset

```
In [39]: # find null counts, percentage of null values, and column type
null_count = housing_testing_data.isnull().sum()
null_percentage = housing_testing_data.isnull().sum() * 100 / len(housing_testing_data)
column_type = housing_testing_data.dtypes

# show null counts, percentage of null values, and column type for columns with more than 1 null
null_summary = pd.concat([null_count, null_percentage, column_type], axis=1, keys=['Missing Count', 'Missing Percentage', 'Data Type'])
null_summary_only_missing = null_summary[null_count != 0].sort_values('Missing Percentage', ascending=False)
null_summary_only_missing
```

Out[39]:

	Missing Count	Percentage Missing	Column Type
PoolQC	1456	99.794380	object
MiscFeature	1408	96.504455	object
Alley	1352	92.666210	object
Fence	1169	80.123372	object
FireplaceQu	730	50.034270	object
LotFrontage	227	15.558602	float64
GarageCond	78	5.346127	object
GarageYrBlt	78	5.346127	float64
GarageQual	78	5.346127	object
GarageFinish	78	5.346127	object
GarageType	76	5.209047	object
BsmtCond	45	3.084304	object
BsmtExposure	44	3.015764	object
BsmtQual	44	3.015764	object
BsmtFinType1	42	2.878684	object
BsmtFinType2	42	2.878684	object
MasVnrType	16	1.096642	object
MasVnrArea	15	1.028101	float64
MSZoning	4	0.274160	object
BsmtFullBath	2	0.137080	float64
BsmtHalfBath	2	0.137080	float64
Functional	2	0.137080	object
Utilities	2	0.137080	object
GarageCars	1	0.068540	float64
GarageArea	1	0.068540	float64
TotalBsmtSF	1	0.068540	float64
KitchenQual	1	0.068540	object
BsmtUnfSF	1	0.068540	float64
BsmtFinSF2	1	0.068540	float64
BsmtFinSF1	1	0.068540	float64
Exterior2nd	1	0.068540	object
Exterior1st	1	0.068540	object
SaleType	1	0.068540	object

```
In [40]: # PoolQC, MiscFeature, Alley, Fence all have over 50% of missing values, we will remove them
housing_testing_data.drop(['Alley', 'PoolQC', 'Fence', 'MiscFeature'], axis=1, inplace=True)

columns_None = ['SaleType', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'GarageType']
# set Nulls in non-numeric columns to 'None'
housing_testing_data[columns_None] = housing_testing_data[columns_None].fillna('None')
```

```
In [41]: # change Null values to 0 for the following variables
columns_zero = ['MasVnrArea', 'GarageArea', 'GarageCars', 'TotalBsmtSF', 'BsmtUnfSF', 'Bsmt
housing_testing_data[columns_zero] = housing_testing_data[columns_zero].fillna(0)

# fill Nulls for Lot Frontage with median value
housing_testing_data['LotFrontage'].fillna(housing_testing_data['LotFrontage'].median()

# fill Nulls with year garage was built with median value
housing_testing_data['GarageYrBlt'].fillna(housing_testing_data['GarageYrBlt'].median()
```

```
In [42]: # convert Paved Drive to dichotomous, indicator variable
housing_testing_data['PavedDrive'] = np.where(housing_testing_data['PavedDrive'] == '\
housing_testing_data['PavedDrive'].value_counts()

# dummy encode the Street, Central Air, Paved Drive indicator variables, we will exclu
housing_testing_data = pd.get_dummies(housing_testing_data, columns=['Street', 'Central
```

```
Out[42]: Y    1301
         N     158
         Name: PavedDrive, dtype: int64
```

```
In [43]: # for heatingQC encode this to be a binary variable to match how we encoded this column
housing_testing_data['HeatingEx'] = np.where(housing_testing_data['HeatingQC'] == 'Ex'
housing_testing_data.drop(columns=['HeatingQC'], inplace=True)
```

```
In [44]: # encode categorical columns with ordinal values
important_categorical = ['ExterQual', 'BsmtQual', 'FireplaceQu', 'KitchenQual', 'GarageQual']

ordinal_mapping = {
    'Ex': 4,
    'Gd': 3,
    'TA': 2,
    'Fa': 1,
    'None': 1,
    'Po': 0,
}

# process columns, replace to categorical features with ordinal ranking
for i in important_categorical:
    housing_testing_data[i] = housing_testing_data[i].replace(ordinal_mapping)

# shape
print('Shape all_data: {}'.format(housing_testing_data.shape))

Shape all data: (1459, 76)
```

```
In [45]: # create new variable TotalSF
housing_testing_data['TotalSF'] = housing_testing_data['TotalBsmtSF'] + housing_testing_data['AboveGrndArea']
```



```
# create new variable, years since the house has been remodeled from selling date (use
housing_testing_data['YrSinceRemod'] = housing_testing_data['YrSold'] - housing_testir
```

In [46]: `### Dummy Encoding`

```
In [47]: # create new df with current numeric columns
housing_testing_numeric_data = pd.DataFrame(housing_testing_data.select_dtypes(exclude=

# create new df with current categorical columns
housing_testing_object_data = pd.DataFrame(housing_testing_data.select_dtypes(exclude=
orig_object_cols = housing_testing_object_data.columns

# dummy encode the categorical columns
housing_testing_data_dummy = pd.get_dummies(housing_testing_object_data, columns=orig_

# create new df with original numeric and new dummy encoded columns
housing_testing_data_large = pd.concat([housing_testing_numeric_data, housing_testing_
housing_testing_data_large
```

Out[47]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1461	20	80.0	11622	5	6	1961	1961
1	1462	20	81.0	14267	6	6	1958	1958
2	1463	60	74.0	13830	5	5	1997	1998
3	1464	60	78.0	9978	6	6	1998	1998
4	1465	120	43.0	5005	8	5	1992	1992
...
1454	2915	160	21.0	1936	4	7	1970	1970
1455	2916	160	21.0	1894	4	5	1970	1970
1456	2917	20	160.0	20000	5	7	1960	1996
1457	2918	85	62.0	10441	5	5	1992	1992
1458	2919	60	74.0	9627	7	5	1993	1994

1459 rows × 246 columns

In [48]: `# create new dataframe with correlated variables with SalePrice observed in the traini`
`feature_list = feature_list.drop('SalePrice')`
`housing_testing_data_large_subset = housing_testing_data_large[feature_list]`
`housing_testing_data_large_subset`

Out[48]:

	LotFrontage	LotArea	OverallQual	YearBuilt	YearRemodAdd	MasVnrArea	ExterQual	BsmtQual
0	80.0	11622	5	1961	1961	0.0	2	
1	81.0	14267	6	1958	1958	108.0	2	
2	74.0	13830	5	1997	1998	0.0	2	
3	78.0	9978	6	1998	1998	20.0	2	
4	43.0	5005	8	1992	1992	0.0	3	
...
1454	21.0	1936	4	1970	1970	0.0	2	
1455	21.0	1894	4	1970	1970	0.0	2	
1456	160.0	20000	5	1960	1996	0.0	2	
1457	62.0	10441	5	1992	1992	0.0	2	
1458	74.0	9627	7	1993	1994	94.0	2	

1459 rows × 50 columns

```

In [49]: # Find the common columns
common_columns = housing_testing_data_large.columns.intersection(housing_training_data_large.columns)

# Keep only the common columns in both DataFrames
housing_testing_data_large_common = housing_testing_data_large[common_columns]
housing_training_data_large_common = housing_training_data_large[common_columns]

In [50]: housing_testing_data_large_common
housing_training_data_large_common
common_columns

```

Out[50]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1461	20	80.0	11622	5	6	1961	1961
1	1462	20	81.0	14267	6	6	1958	1958
2	1463	60	74.0	13830	5	5	1997	1998
3	1464	60	78.0	9978	6	6	1998	1998
4	1465	120	43.0	5005	8	5	1992	1992
...
1454	2915	160	21.0	1936	4	7	1970	1970
1455	2916	160	21.0	1894	4	5	1970	1970
1456	2917	20	160.0	20000	5	7	1960	1996
1457	2918	85	62.0	10441	5	5	1992	1992
1458	2919	60	74.0	9627	7	5	1993	1994

1459 rows × 240 columns

Out[50]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1	60	65.0	8450	7	5	2003	2003
1	2	20	80.0	9600	6	8	1976	1976
2	3	60	68.0	11250	7	5	2001	2002
3	4	70	60.0	9550	7	5	1915	1970
4	5	60	84.0	14260	8	5	2000	2000
...
1455	1456	60	62.0	7917	6	5	1999	2000
1456	1457	20	85.0	13175	6	6	1978	1988
1457	1458	70	66.0	9042	7	9	1941	2006
1458	1459	20	68.0	9717	5	6	1950	1996
1459	1460	20	75.0	9937	5	6	1965	1965

1455 rows × 240 columns

```
Out[50]: Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
              'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'ExterQual',
              ...,
              'SaleType_ConLw', 'SaleType_New', 'SaleType_Oth', 'SaleType_WD',
              'SaleCondition_Abnorml', 'SaleCondition_AdjLand',
              'SaleCondition_Alloca', 'SaleCondition_Family', 'SaleCondition_Normal',
              'SaleCondition_Partial'],
              dtype='object', length=240)
```

In []:

Constructing Models to Predict Home Prices

Model Assumptions

1. Linearity
2. Homoscedasticity
3. Independence of Errors
4. Multivariate Normality
5. No or little Multicollinearity

Below are simple and multiple regressions that examine the associations between variables of interest and sale price.

```
In [51]: import numpy as np
import statsmodels.api as sm
# New feature is highly correlated, lets try a simple linear regression
x = housing_training_data['TotalSF']
y = housing_training_data['SalePrice']

#add constant to predictor variables
X = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y, X).fit()

#view model summary
print(model.summary())

# plot the regression model
sns.regplot(x=x, y=y)
```

OLS Regression Results

```

=====
Dep. Variable:          SalePrice      R-squared:                0.672
Model:                  OLS           Adj. R-squared:           0.672
Method:                 Least Squares  F-statistic:              2982.
Date:                   Sun, 16 Apr 2023  Prob (F-statistic):      0.00
Time:                   16:12:31       Log-Likelihood:           -17613.
No. Observations:      1455          AIC:                     3.523e+04
Df Residuals:          1453          BIC:                     3.524e+04
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-3.239e+04	4054.647	-7.989	0.000	-4.03e+04	-2.44e+04
TotalSF	83.1361	1.522	54.610	0.000	80.150	86.122

```

=====
Omnibus:                125.366      Durbin-Watson:           1.967
Prob(Omnibus):           0.000      Jarque-Bera (JB):        649.680
Skew:                    0.197      Prob(JB):                8.39e-142
Kurtosis:                6.250      Cond. No.                9.41e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.41e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Out[51]: <AxesSubplot:xlabel='TotalSF', ylabel='SalePrice'>




```

In [52]: # plot the residuals
y_pred=model.predict(X)
sns.residplot(x=y, y=y_pred)

```

Out[52]: <AxesSubplot:xlabel='SalePrice'>



The residual plot shows evidence of Heteroscedasticity since the residuals are not evenly scattered. For higher Sales Prices, the residuals are negative indicating that the model is over estimating homes with higher Sales Prices. There is evidence that this model violates the linearity assumption, Homoscedasticity assumption and the Independence of Errors assumption.


```
In [53]: # qqplot
import matplotlib.pyplot as plt
stats.probplot(y-y_pred, dist="norm", plot=plt)
plt.title("Normal QQ Plot")
plt.xlabel("Theoretical Quantiles")
plt.ylabel("Sample Quantiles")
plt.show()
```

```
Out[53]: ((array([-3.30417817, -3.04690148, -2.90382339, ..., 2.90382339,
3.04690148, 3.30417817]),
array([-184865.74057368, -169571.36528068, -156517.81407152, ...,
195404.36999739, 236556.60682101, 253807.82492125])),
(42839.88750440423, 1.5086045390486717e-10, 0.9775030165443651))
```

```
Out[53]: Text(0.5, 1.0, 'Normal QQ Plot')
```

```
Out[53]: Text(0.5, 0, 'Theoretical Quantiles')
```

```
Out[53]: Text(0, 0.5, 'Sample Quantiles')
```



The tails of the distribution deviate from the qqline indicating that the errors are not Normally distributed.

Let's try transforming the independent variable, Sales Price since we saw earlier in this analysis that this helped to Normalize its distribution.

```
In [54]: # Log transform Sales Price variable
y_log = np.log(housing_training_data['SalePrice'])

#add constant to predictor variables
X = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y_log, X).fit()

#view model summary
print(model.summary())

# plot the regression model
sns.regplot(x=x, y=y_log)
```

OLS Regression Results

```
=====
Dep. Variable:          SalePrice    R-squared:                0.669
Model:                  OLS          Adj. R-squared:           0.669
Method:                 Least Squares  F-statistic:             2934.
Date:                  Sun, 16 Apr 2023  Prob (F-statistic):       0.00
Time:                  16:12:32        Log-Likelihood:          89.766
No. Observations:      1455           AIC:                    -175.5
Df Residuals:          1453           BIC:                    -165.0
Df Model:               1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	10.9256	0.021	518.068	0.000	10.884	10.967
TotalSF	0.0004	7.92e-06	54.167	0.000	0.000	0.000


```
=====
Omnibus:                274.748    Durbin-Watson:           1.937
Prob(Omnibus):           0.000     Jarque-Bera (JB):        567.536
Skew:                   -1.090     Prob(JB):                5.77e-124
Kurtosis:                5.147     Cond. No.                 9.41e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.41e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
Out[54]: <AxesSubplot:xlabel='TotalSF', ylabel='SalePrice'>
```




The relationship between $\log(\text{Sale Price})$ and TotalSF appears to be linear. This meets the Linearity assumption.

An r-squared value of 0.669 means that the model explains 66.9 percent of the variance in the dependent variable. The omnibus test, however, indicates that residuals are not normally distributed and the high kurtosis value indicates that the distribution of the residuals are more peaked than a normal distribution. The condition number is 1.12×10^4 , which is quite high and suggests that there may be multicollinearity in the model.

```
In [55]: # plot the residuals
y_pred=model.predict(X)
sns.residplot(x=y_log, y=y_pred)
```

```
Out[55]: <AxesSubplot:xlabel='SalePrice'>
```




The residuals appear to be more randomly scattered across values of Sales Price. This appears to better meet the Homoscedasticity assumption and the Independence of Errors assumptions than the model with the untransformed Sales Price.

```
In [56]: # qqplot
stats.probplot(y_log-y_pred, dist="norm", plot=plt)
plt.title("Normal QQ Plot")
plt.xlabel("Theoretical Quantiles")
```



```
plt.ylabel("Sample Quantiles")
plt.show()
```

```
Out[56]: ((array([-3.30417817, -3.04690148, -2.90382339, ...,  2.90382339,
        3.04690148,  3.30417817]),
        array([-1.17224411, -1.08303201, -1.05546774, ...,  0.54315541,
        0.55983944,  0.73613196])),
        (0.22111127940895553, 6.253817917920981e-15, 0.9700026442288759))
Out[56]: Text(0.5, 1.0, 'Normal QQ Plot')
Out[56]: Text(0.5, 0, 'Theoretical Quantiles')
Out[56]: Text(0, 0.5, 'Sample Quantiles')
```



The tails of the distribution still stray from the qqline, indicating the distribution of the errors is not Normal.

We can see the assumptions of a linear model are still not met by transforming Sales Price, lets try building a polynomial model.

Polynomial Regression

Third Order Polynomial Regression with Total SF as a predictor of Sales Price

```
In [57]: from sklearn.preprocessing import PolynomialFeatures
x = housing_training_data[['TotalSF']]
y = housing_training_data['SalePrice']

polynomial_features= PolynomialFeatures(degree=3)
xp = polynomial_features.fit_transform(x)
xp.shape

model = sm.OLS(y, xp).fit()

#view model summary
print(model.summary())
# predicted sales price
y_pred = model.predict(xp)

# plot model against data
plt.scatter(x,y)
plt.plot(x,y_pred)
```

Out[57]: (1455, 4)

```

OLS Regression Results
=====
Dep. Variable:          SalePrice    R-squared:                0.689
Model:                  OLS          Adj. R-squared:           0.688
Method:                 Least Squares  F-statistic:             1070.
Date:                  Sun, 16 Apr 2023  Prob (F-statistic):       0.00
Time:                  16:12:34       Log-Likelihood:          -17576.
No. Observations:      1455          AIC:                    3.516e+04
Df Residuals:          1451          BIC:                    3.518e+04
Df Model:               3
Covariance Type:       nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const      6.831e+04   2.17e+04     3.146     0.002    2.57e+04   1.11e+05
x1         -4.1538     24.479    -0.170     0.865    -52.171    43.864
x2          0.0203      0.009     2.311     0.021     0.003     0.037
x3        -1.007e-06   9.93e-07    -1.015     0.310    -2.95e-06   9.4e-07
=====
Omnibus:             126.279   Durbin-Watson:           1.954
Prob(Omnibus):        0.000   Jarque-Bera (JB):        633.593
Skew:                 -0.222   Prob(JB):                 2.61e-138
Kurtosis:              6.202   Cond. No.                 5.70e+11
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.7e+11. This might indicate that there are strong multicollinearity or other numerical problems.

Out[57]: <matplotlib.collections.PathCollection at 0x7fee806f4950>

Out[57]: [<matplotlib.lines.Line2D at 0x7fee8085ae10>]




```

In [58]: # plot the residuals
sns.residplot(x=y, y=y_pred)

```

Out[58]: <AxesSubplot:xlabel='SalePrice'>



The residuals are fairly scattered across Sales Prices. Let's try adding more predictors to a linear regression model.

Multiple Linear Regression

Check the correlation between the two new variables. If they are highly correlated we won't construct a multiple linear regression model with both of those variables as predictors.

```
In [59]: df_corr_mult_lreg = housing_training_data[['TotalSF', 'YrSinceRemod']]
df_corr_mult_lreg.corr()
```

```
Out[59]:
```

	TotalSF	YrSinceRemod
TotalSF	1.000000	-0.352279
YrSinceRemod	-0.352279	1.000000

```
In [ ]:
```

TotalSF and YrSinceRemod are not highly correlated so we will construct a multiple linear regression model using the two variables as predictors of the log transformed Sales Price.

```
In [60]: x = housing_training_data[['TotalSF', 'YrSinceRemod']]
y_log = np.log(housing_training_data['SalePrice'])

#add constant to predictor variables
X = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y_log, X).fit()

#view model summary
print(model.summary())
```

OLS Regression Results

=====						
Dep. Variable:	SalePrice	R-squared:	0.761			
Model:	OLS	Adj. R-squared:	0.760			
Method:	Least Squares	F-statistic:	2308.			
Date:	Sun, 16 Apr 2023	Prob (F-statistic):	0.00			
Time:	16:12:34	Log-Likelihood:	326.20			
No. Observations:	1455	AIC:	-646.4			
Df Residuals:	1452	BIC:	-630.6			
Df Model:	2					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	11.2212	0.022	513.127	0.000	11.178	11.264
TotalSF	0.0004	7.19e-06	51.300	0.000	0.000	0.000
YrSinceRemod	-0.0062	0.000	-23.614	0.000	-0.007	-0.006
=====						
Omnibus:	279.765	Durbin-Watson:	1.927			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	639.242			
Skew:	-1.065	Prob(JB):	1.55e-139			
Kurtosis:	5.450	Cond. No.	1.15e+04			
=====						

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.15e+04. This might indicate that there are strong multicollinearity or other numerical problems.

An r-squared value of 0.765 means that the model explains 76.5 percent of the variance in the dependent variable. The adjusted R-squared value is about the same as the r-squared value, indicating that we aren't overfitting the model by adding multiple variables. The omnibus test, however, indicates that residuals are not normally distributed and the high kurtosis value indicates that the distribution of the residuals are more peaked than a normal distribution. The condition number is 1.12e+04, which is quite high and suggests that there may be multicollinearity in the model.

```
In [61]: # plot the residuals
y_pred=model.predict(X)
sns.residplot(x=y_log, y=y_pred)
```

```
Out[61]: <AxesSubplot:xlabel='SalePrice'>
```



The residuals appear to be more randomly scattered across values of Sales Price. This appears to better meet the Homoscedasticity assumption and the Independence of Errors assumptions than the model with the untransformed Sales Price. But as we mentioned above, the high omnibus value indicates the distribution of the residuals is not normally distributed.

Piecewise Regression

We can try fitting a linear regression model of sale price using a piecewise regression model.

```
In [62]: import pwlf

x = np.array(housing_training_data['TotalSF'])
y = np.array(np.log(housing_training_data['SalePrice']))

# initialize piecewise linear fit with your x and y data
my_pwlf = pwlf.PiecewiseLinFit(x, y)

# fit the data for four line segments
res = my_pwlf.fit(3)

# predict for the determined points
xHat = np.array(housing_training_data['TotalSF'])
yHat = my_pwlf.predict(xHat)

piecewise_regression_output = housing_training_data[["TotalSF"]]
piecewise_regression_output['Log Sale Price'] = y.tolist()
piecewise_regression_output['Predicted Log Sale Price'] = yHat.tolist()
piecewise_regression_output['residual'] = piecewise_regression_output['Log Sale Price'] - yHat.tolist()
piecewise_regression_output['squared residuals'] = piecewise_regression_output['residual']**2

RMSE = (piecewise_regression_output['squared residuals'].sum() / len(piecewise_regression_output))
print(f"The Root Mean Squared Error of this piecewise regression model is {RMSE}.")

correlation = piecewise_regression_output['Log Sale Price'].corr(piecewise_regression_output['Predicted Log Sale Price'])
print(f"The correlation of this piecewise regression model is {correlation}.")
```

The Root Mean Squared Error of this piecewise regression model is 0.22538513350191516.

The correlation of this piecewise regression model is 0.8215282922354664.

```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

We can visualize the residuals plot to examine whether the regression model may violate any of the regression assumptions.

```


In [63]: sns.residplot(x=y, y=yHat)
plt.ylabel('Standardized Residual')
plt.xlabel('TotalSF')

```

Out[63]: <AxesSubplot:>

Out[63]: Text(0, 0.5, 'Standardized Residual')


Out[63]: Text(0.5, 0, 'TotalSF')



We can visualize a Q-Q plot to determine whether this regression model violates the linear regression assumption that residuals are normally distributed.

```
In [64]: # qqplot
stats.probplot(y - yHat, dist="norm", plot=plt)
plt.title("Normal QQ Plot")
plt.xlabel("Theoretical Quantiles")
plt.ylabel("Sample Quantiles")
plt.show()

Out[64]: ((array([-3.30417817, -3.04690148, -2.90382339, ...,  2.90382339,
        3.04690148,  3.30417817]),
        array([-1.1651156 , -1.0640493 , -1.0393697 , ...,  0.54546573,
        0.57744714,  0.72360866])),
        (0.2193429300027091, 3.1475258810892375e-15, 0.9712466424655569))
Out[64]: Text(0.5, 1.0, 'Normal QQ Plot')
Out[64]: Text(0.5, 0, 'Theoretical Quantiles')
Out[64]: Text(0, 0.5, 'Sample Quantiles')
```



We will try and add more predictors to our model. First, let's observe any multicollinearity.

Inspection of multicollinearity: VIF, correlations

The correlation between garage cars and total square feet is moderately high. We will take that into consideration when analyzing the output of the model.

```
In [65]: from statsmodels.stats.outliers_influence import variance_inflation_factor

# the independent variables set
x = housing_training_data[['TotalSF', 'YrSinceRemod', 'GarageCars', 'ExterQual', 'CentralAir']]

# VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = x.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(x.values, i)
                   for i in range(len(x.columns))]
```

```
print(vif_data)
x.corr()
```

	feature	VIF
0	TotalSF	19.799177
1	YrSinceRemod	1.901546
2	GarageCars	11.051476
3	ExterQual	22.737370
4	CentralAir_Y	12.377024

Out[65]:

	TotalSF	YrSinceRemod	GarageCars	ExterQual	CentralAir_Y
TotalSF	1.000000	-0.352279	0.556693	0.528016	0.180174
YrSinceRemod	-0.352279	1.000000	-0.422033	-0.587649	-0.299245
GarageCars	0.556693	-0.422033	1.000000	0.524166	0.233414
ExterQual	0.528016	-0.587649	0.524166	1.000000	0.206058
CentralAir_Y	0.180174	-0.299245	0.233414	0.206058	1.000000

These VIF values suggests that our model contains some multicollinearity across all features.

Some resources stated that greater than 10 suggests multicollinearity, while others used a cut-off of 5.0 or 1.0. This may not be an issue for a predictive model, but could have implications for an inferential model.

```
In [66]: x = housing_training_data[['TotalSF', 'YrSinceRemod', 'GarageCars', 'ExterQual', 'CentralAir_Y']]
y_log = np.log(housing_training_data['SalePrice'])
```

```
polynomial_features = PolynomialFeatures(degree=2)
xp = polynomial_features.fit_transform(x)
xp.shape
```

```
#add constant to predictor variables
x = sm.add_constant(xp)
```

```
#fit polynomial regression model
model = sm.OLS(y_log, x).fit()
```

```
#view model summary
print(model.summary())
```

Out[66]: (1455, 21)

OLS Regression Results

=====						
Dep. Variable:	SalePrice	R-squared:	0.838			
Model:	OLS	Adj. R-squared:	0.835			
Method:	Least Squares	F-statistic:	389.4			
Date:	Sun, 16 Apr 2023	Prob (F-statistic):	0.00			
Time:	16:12:36	Log-Likelihood:	608.02			
No. Observations:	1455	AIC:	-1176.			
Df Residuals:	1435	BIC:	-1070.			
Df Model:	19					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	10.7869	0.166	65.166	0.000	10.462	11.112
x1	0.0004	5.06e-05	8.017	0.000	0.000	0.001
x2	-0.0022	0.003	-0.826	0.409	-0.007	0.003
x3	0.1521	0.051	2.955	0.003	0.051	0.253
x4	-0.0450	0.096	-0.468	0.640	-0.234	0.144
x5	0.0780	0.065	1.195	0.232	-0.050	0.206
x6	-2.858e-08	6.91e-09	-4.135	0.000	-4.21e-08	-1.5e-08
x7	-1.603e-06	4.55e-07	-3.523	0.000	-2.5e-06	-7.11e-07
x8	3.71e-06	1.1e-05	0.336	0.737	-1.8e-05	2.54e-05
x9	3.17e-05	1.51e-05	2.100	0.036	2.09e-06	6.13e-05
x10	-1.362e-05	2.64e-05	-0.516	0.606	-6.55e-05	3.82e-05
x11	-1.634e-05	1.74e-05	-0.939	0.348	-5.05e-05	1.78e-05
x12	-0.0011	0.000	-2.609	0.009	-0.002	-0.000
x13	0.0017	0.001	1.910	0.056	-4.62e-05	0.003
x14	0.0018	0.001	1.812	0.070	-0.000	0.004
x15	-0.0348	0.008	-4.562	0.000	-0.050	-0.020
x16	0.0361	0.019	1.870	0.062	-0.002	0.074
x17	-0.0044	0.023	-0.193	0.847	-0.049	0.040
x18	-0.0057	0.018	-0.320	0.749	-0.040	0.029
x19	-0.0054	0.058	-0.092	0.926	-0.119	0.108
x20	0.0780	0.065	1.195	0.232	-0.050	0.206
=====						
Omnibus:	240.889	Durbin-Watson:	1.985			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	706.892			
Skew:	-0.845	Prob(JB):	3.17e-154			
Kurtosis:	5.967	Cond. No.	1.01e+21			
=====						


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 9.79e-26. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [67]: # plot the residuals
y_pred=model.predict(x)
sns.residplot(x=y_log, y=y_pred)
```

```
Out[67]: <AxesSubplot:xlabel='SalePrice'>
```



The adjusted r-squared value is high for the 2nd order polynomial model, but the residuals are not randomly scattered. The plot indicates that the model is underestimating the value of homes with actual low sales prices and overestimating the value of homes with actual high sales prices.

```
In [68]: from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression
from numpy import mean
from numpy import absolute
from numpy import sqrt
from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
from sklearn.metrics import mean_squared_error
```

Let's use PCA to find some more important predictors to add to our models

Regress on principal components

```
In [69]: from sklearn.decomposition import PCA
from sklearn.preprocessing import scale
from sklearn.linear_model import LinearRegression
import os
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score

scaler = StandardScaler()

pca = PCA(n_components=8)

## independent variables ###

x_raw = housing_training_data.select_dtypes(exclude=['object']).drop(columns = ['SalePrice'])
x_scale = scaler.fit_transform(x_raw)

x_pca_raw = pca.fit_transform(x_raw)
x_pca_scale = pca.fit_transform(x_scale) # PCA is affected by scale. We use the scaled
```

```

## importance scores of the principal components ##
# show loadings for features
loadings = pd.DataFrame(pca.components_.T, columns=['PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6', 'PC7', 'PC8', 'PC9', 'PC10'])
loadings

## dependent variable ##

y_trans = scaler.fit_transform(housing_training_data[['SalePrice']])

y_raw = np.array(housing_training_data[['SalePrice']]).reshape(-1,1)
y_scale = np.array(y_trans).reshape(-1,1)

## train linear model ##

regr = LinearRegression()

regr.fit(x_pca_scale, y_scale)

y_pred = regr.predict(x_pca_scale)

plt.subplot(1, 2, 1)

plt.scatter(y_scale, y_pred)
plt.xlabel('Log Sale Price (Standard Scaled)')
plt.ylabel('Predicted Log Sale Price')
plt.title('Log Sale Price vs Predicted Log Sale Price')

# Residuals
plt.subplot(1, 2, 2)
sns.residplot(x=y_scale, y=y_pred)
plt.xlabel('Log Sale Price')
plt.ylabel('Residual')
plt.title('Residual Plot')
## calculate RMSE ##

# Mean Squared Error
MSE = np.square(np.subtract(y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_scale, y_pred)
print("R_sq:",r2)

```

Out[69]:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
Id	0.003671	0.007609	-0.012863	-0.018848	-0.031304	0.014721	-0.000231	-0.230336
MSSubClass	0.013197	0.056123	-0.258946	-0.164299	0.272634	0.191425	-0.164072	-0.186888
LotFrontage	-0.098662	0.109474	0.215217	0.061892	-0.136464	-0.021591	0.195174	0.212405
LotArea	-0.060616	0.085147	0.241554	-0.042553	0.031626	-0.034394	0.141926	0.328845
OverallQual	-0.256954	0.013253	-0.057438	0.035494	0.024299	-0.095312	-0.106163	-0.076213
OverallCond	0.055815	0.014262	0.008466	-0.135176	-0.040599	-0.481640	0.270476	-0.152957
YearBuilt	-0.219252	-0.220127	-0.134601	0.005329	0.012171	0.181039	-0.024807	0.179010
YearRemodAdd	-0.202447	-0.135251	-0.224474	0.081358	0.105161	-0.206615	0.237536	0.013918
MasVnrArea	-0.143366	0.042596	0.050969	-0.027038	0.048708	0.116316	-0.144984	-0.030711
ExterQual	-0.233778	-0.064617	-0.114988	0.116459	0.063804	-0.076014	-0.067799	-0.095841
BsmtQual	-0.229908	-0.114895	-0.085267	0.067084	0.123963	0.021872	-0.090064	0.080774
BsmtFinSF1	-0.099578	-0.153913	0.299269	-0.151573	0.374517	0.060625	0.034273	-0.065014
BsmtFinSF2	0.008778	-0.017667	0.170686	-0.097427	0.022247	-0.067193	0.181676	0.085114
BsmtUnfSF	-0.096454	0.126643	-0.099844	0.415105	-0.359291	-0.007090	-0.097226	-0.009365
TotalBsmtSF	-0.205445	-0.032307	0.274765	0.249944	0.015222	0.029578	0.003010	-0.044868
1stFlrSF	-0.191291	0.047038	0.322391	0.228092	-0.012687	0.051398	0.047960	-0.094858
2ndFlrSF	-0.083826	0.325601	-0.263244	-0.286736	0.078471	-0.039820	-0.004822	0.045017
LowQualFinSF	0.016325	0.121673	0.010900	0.022469	0.018783	-0.096268	0.098606	-0.087547
GrLivArea	-0.212725	0.330630	0.009550	-0.079419	0.060909	-0.006249	0.040909	-0.039396
BsmtFullBath	-0.061426	-0.165104	0.243546	-0.121882	0.422020	0.067214	0.096279	-0.103952
BsmtHalfBath	0.012224	-0.002542	0.046869	-0.089462	-0.056439	-0.142792	0.079802	0.210566
FullBath	-0.201502	0.150867	-0.134477	0.064371	0.018548	0.139731	0.098124	0.012173
HalfBath	-0.093650	0.134351	-0.197361	-0.319660	0.076294	-0.018475	-0.112438	0.178204
BedroomAbvGr	-0.054909	0.356909	-0.017353	-0.089966	-0.083644	0.073435	0.242245	0.114330
KitchenAbvGr	0.041675	0.189837	-0.014183	0.086211	0.119344	0.377492	0.156014	-0.164956
KitchenQual	-0.224767	-0.072831	-0.089469	0.069620	0.071262	-0.165911	0.015015	-0.123078
TotRmsAbvGrd	-0.160080	0.368444	-0.015952	-0.048376	0.017501	0.055540	0.135361	-0.016877
Fireplaces	-0.140927	0.115976	0.199048	-0.128389	0.035761	-0.219753	-0.387406	0.022184
FireplaceQu	-0.162657	0.098898	0.121926	-0.017785	0.001883	-0.234443	-0.425147	-0.046188
GarageYrBltd	-0.217834	-0.202084	-0.162327	-0.010719	-0.048295	0.184504	0.047086	0.077313
GarageCars	-0.235138	-0.029470	0.027626	-0.079300	-0.147468	0.178528	0.004658	-0.089752
GarageArea	-0.224938	-0.039640	0.079160	-0.078222	-0.156884	0.177333	0.047977	-0.111435
GarageQual	-0.123336	-0.120531	0.066392	-0.325948	-0.343787	0.082770	0.031902	-0.204625
GarageCond	-0.121213	-0.139310	0.070769	-0.330448	-0.357112	0.073309	0.032527	-0.173778

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
WoodDeckSF	-0.104889	-0.019359	0.068648	-0.091644	0.117418	-0.031780	0.130396	0.072254
OpenPorchSF	-0.109004	0.064471	-0.040801	0.012955	0.060412	-0.076822	-0.057477	0.156020
EnclosedPorch	0.070821	0.133254	0.054251	0.045052	0.002906	-0.130765	0.058703	-0.383462
3SsnPorch	-0.016219	-0.018176	0.021199	0.031480	-0.041260	-0.060069	0.052732	0.065086
ScreenPorch	-0.024086	0.050690	0.111115	-0.107502	-0.031379	-0.128971	-0.230816	0.025060
PoolArea	-0.008178	0.038436	0.050052	-0.086850	-0.021028	-0.046144	0.093498	-0.148385
MiscVal	0.011021	0.023471	0.012395	-0.031549	0.006030	-0.034770	0.141798	-0.014292
MoSold	-0.019893	0.031797	0.001362	0.020580	-0.024677	-0.040583	-0.106347	0.175026
YrSold	0.007311	-0.029935	0.015310	-0.009106	0.065255	0.002643	0.143420	-0.129127
Street_Pave	-0.012603	0.006765	-0.080769	0.021471	-0.031430	-0.114273	0.005132	-0.366851
CentralAir_Y	-0.112814	-0.148689	0.011083	-0.165249	-0.133699	-0.179586	0.102746	0.064415
PavedDrive_Y	-0.108617	-0.164614	0.035119	-0.175303	-0.145556	0.098251	0.001568	0.098963
HeatingEx	-0.165890	-0.093989	-0.152127	0.112837	0.039147	-0.138834	0.035621	-0.020257
YrSinceRemod	0.202978	0.133363	0.225526	-0.081969	-0.100990	0.206846	-0.228373	-0.022236
TotalSF	-0.250927	0.199534	0.155107	0.083260	0.048228	0.011921	0.028484	-0.050165

Out[69]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Out[69]: <AxesSubplot:>

Out[69]: <matplotlib.collections.PathCollection at 0x7fee82082650>

Out[69]: Text(0.5, 0, 'Log Sale Price (Standard Scaled)')

Out[69]: Text(0, 0.5, 'Predicted Log Sale Price')

Out[69]: Text(0.5, 1.0, 'Log Sale Price vs Predicted Log Sale Price')

Out[69]: <AxesSubplot:>

Out[69]: <AxesSubplot:>

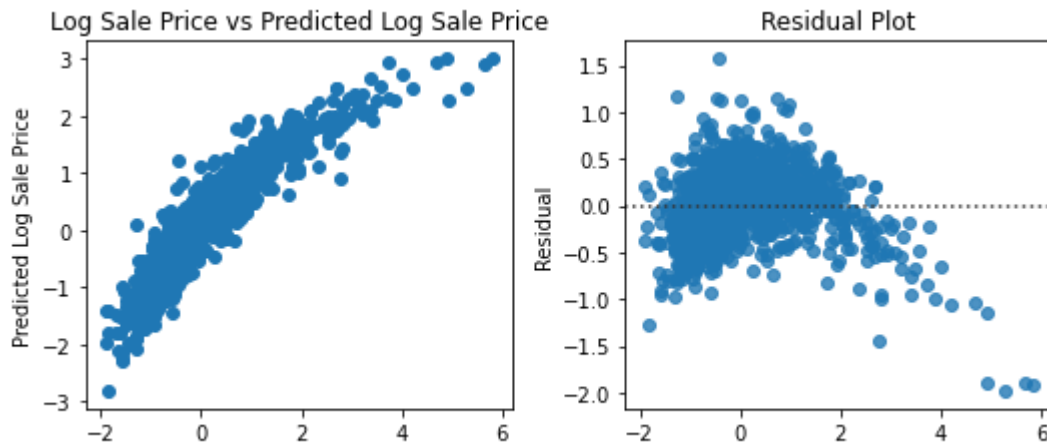
Out[69]: Text(0.5, 0, 'Log Sale Price')

Out[69]: Text(0, 0.5, 'Residual')

Out[69]: Text(0.5, 1.0, 'Residual Plot')

MSE: 0.15473155070155967

R_sq: 0.8452684492984404



In [70]: `x_raw.columns`

Out[70]: Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'ExterQual',
'BsmtQual', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
'1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath',
'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
'KitchenQual', 'TotRmsAbvGrd', 'Fireplaces', 'FireplaceQu',
'GarageYrBlt', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'Street_Pave',
'CentralAir_Y', 'PavedDrive_Y', 'HeatingEx', 'YrSinceRemod', 'TotalSF'],
dtype='object')

```
In [71]: #define cross-validation method to use

cv = KFold(n_splits=10, random_state=1, shuffle=True)

#build multiple linear regression model
model = LinearRegression()

#use k-fold CV to evaluate model
scores_mse = cross_val_score(model, x_pca_scale, y_scale, scoring='neg_mean_absolute_e
                             cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f MSE with a standard deviation of %0.2f" % (scores_mse.mean(), scores_mse.std()))

#use k-fold CV to evaluate model R2
scores_r2 = cross_val_score(model, x_pca_scale, y_scale, scoring='r2',
                             cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f r2 with a standard deviation of %0.2f" % (scores_r2.mean(), scores_r2.std()))

Over 10 folds: -0.28 MSE with a standard deviation of 0.04
Over 10 folds: 0.85 r2 with a standard deviation of 0.03
```

There appears to be a lot of large negative residuals for higher Sales Prices indicating that the model is overestimating the price of homes with large sales prices.

```
In [72]: # Try Log of Y

y_log_scale = scaler.fit_transform(np.array(np.log(housing_training_data['SalePrice'])))

## train linear model ##
```

```

regr2 = LinearRegression()

regr2.fit(x_pca_scale, y_log_scale)

y_pred = regr2.predict(x_pca_scale)

plt.subplot(1, 2, 1)

plt.scatter(y_log_scale, y_pred)
plt.xlabel('Log Sale Price (Standard Scaled)')
plt.ylabel('Predicted Log Sale Price')
plt.title('Log Sale Price vs Predicted Log Sale Price')

# Residuals
plt.subplot(1, 2, 2)
sns.residplot(x=y_log_scale, y=y_pred)
plt.xlabel('Log Sale Price')
plt.ylabel('Residual')
plt.title('Residual Plot')

##RMSE
# Mean Squared Error

MSE = np.square(np.subtract(y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_log_scale, y_pred)
print("R_sq:",r2)

```

Out[72]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Out[72]: <AxesSubplot:>

Out[72]: <matplotlib.collections.PathCollection at 0x7fee820d0c10>

Out[72]: Text(0.5, 0, 'Log Sale Price (Standard Scaled)')

Out[72]: Text(0, 0.5, 'Predicted Log Sale Price')

Out[72]: Text(0.5, 1.0, 'Log Sale Price vs Predicted Log Sale Price')

Out[72]: <AxesSubplot:>


Out[72]: <AxesSubplot:>

Out[72]: Text(0.5, 0, 'Log Sale Price')

Out[72]: Text(0, 0.5, 'Residual')

Out[72]: Text(0.5, 1.0, 'Residual Plot')

MSE: 0.1675065487775354
R_sq: 0.8864767791398497



```
In [73]: #define cross-validation method to use

cv = KFold(n_splits=10, random_state=1, shuffle=True)

#build multiple linear regression model
model = LinearRegression()

#use k-fold CV to evaluate model
scores_mse = cross_val_score(model, x_pca_scale, y_log_scale, scoring='neg_mean_absolute_error',
                              cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f MSE with a standard deviation of %0.2f" % (scores_mse.mean(), scores_mse.std()))

#use k-fold CV to evaluate model R2
scores_r2 = cross_val_score(model, x_pca_scale, y_log_scale, scoring='r2',
                              cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f r2 with a standard deviation of %0.2f" % (scores_r2.mean(), scores_r2.std()))

Over 10 folds: -0.25 MSE with a standard deviation of 0.01
Over 10 folds: 0.88 r2 with a standard deviation of 0.02
```

There are a lot of positive residuals for lower sales prices indicating that the model is underestimating the value of homes with lower sales prices.

```
In [74]: ### FROM PCA ON SALE PRICE WE SEE QUADRATIC PATTERN -- TRY DEGREE 2 POLYNOMIAL

from sklearn.preprocessing import PolynomialFeatures

#define our polynomial model, degree 2

degree=2

# PolynomialFeatures will create a new matrix consisting of all polynomial combinations
# of the features with a degree less than or equal to the degree we just gave the model
poly_model = PolynomialFeatures(degree=degree)

# transform out polynomial features
poly_x_values = poly_model.fit_transform(x_pca_scale)

regression_model = LinearRegression()
regression_model.fit(poly_x_values, y_scale)
y_pred = regression_model.predict(poly_x_values)
```



```

plt.subplot(1, 2, 1)
plt.scatter(y_scale, y_pred)
plt.xlabel('Sale Price (Standard Scaled)')
plt.ylabel('Predicted Sale Price')
plt.title('Sale Price vs Predicted Sale Price')

# Residuals

plt.subplot(1, 2, 2)
sns.residplot(x=y_scale, y=y_pred)
plt.xlabel('Sale Price')
plt.ylabel('Residual')
plt.title('Residual Plot')

# SCORES

SS_Residual = sum((y_scale-y_pred)**2)
SS_Total = sum((y_scale-np.mean(y_scale))**2)
r_squared = 1 - (float(SS_Residual))/SS_Total
adjusted_r_squared = 1 - (1-r_squared)*(len(y_scale)-1)/(len(y)-poly_x_values.shape[1])

print('x_shape_transformed',poly_x_values.shape)
print('x_shape_original',x_pca_scale.shape)

MSE = np.square(np.subtract(y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_scale, y_pred)
print("R_sq:",r_squared[0])

print("R_sq_adjusted:", 1 - (1-r_squared[0])*(len(y_scale)-1)/(len(y_scale)-poly_x_val

```

Out[74]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Out[74]: <AxesSubplot:>

Out[74]: <matplotlib.collections.PathCollection at 0x7fee82271c10>

Out[74]: Text(0.5, 0, 'Sale Price (Standard Scaled)')

Out[74]: Text(0, 0.5, 'Predicted Sale Price')

Out[74]: Text(0.5, 1.0, 'Sale Price vs Predicted Sale Price')

Out[74]: <AxesSubplot:>

Out[74]: <AxesSubplot:>

Out[74]: Text(0.5, 0, 'Sale Price')

Out[74]: Text(0, 0.5, 'Residual')

Out[74]: Text(0.5, 1.0, 'Residual Plot')


x_shape_transformed (1455, 45)

x_shape_original (1455, 8)

MSE: 0.08764742224548526

R_sq: 0.9123525777545145

R_sq_adjusted: 0.9095533343187112



```
In [75]: #define cross-validation method to use

cv = KFold(n_splits=10, random_state=1, shuffle=True)

#build multiple linear regression model
model = LinearRegression()

#use k-fold CV to evaluate model
scores_mse = cross_val_score(model, poly_x_values, y_scale, scoring='neg_mean_absolute',
                             cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f MSE with a standard deviation of %0.2f" % (scores_mse.mean(), scores_mse.std()))

#use k-fold CV to evaluate model R2
scores_r2 = cross_val_score(model, poly_x_values, y_scale, scoring='r2',
                             cv=cv, n_jobs=-1)
print("Over 10 folds: %0.2f r2 with a standard deviation of %0.2f" % (scores_r2.mean(), scores_r2.std()))

Over 10 folds: -0.22 MSE with a standard deviation of 0.02
Over 10 folds: 0.90 r2 with a standard deviation of 0.01
```

We will need to verify that better fit is not just from including more variables. We do this using cross validation to assess out of sample accuracy.

The residuals are a bit more scattered, but there are still some more large negative residuals for higher sales prices.

```
In [76]: # Lastly, try a regularization technique

from sklearn.linear_model import ElasticNet

scaler = StandardScaler()

# As an experiment, use on an original set of x variables...

x_raw = housing_training_data[numerical_vars].drop(columns = 'SalePrice')
x_scale = scale(x_raw)

#x_raw_new = housing_training_data.select_dtypes(exclude=['object']).drop(columns = [
#x_scale = scaler.fit_transform(x_raw_new)

regr = ElasticNet()
```

```

regr.fit(x_scale, y_log_scale)

ElasticNet(random_state=0)
print(regr.coef_)
print(regr.intercept_)

y_pred = regr.predict(x_scale)

plt.subplot(1, 2, 1)
plt.scatter(y_log_scale, y_pred)
plt.xlabel('Sale Price (Standard Scaled)')
plt.ylabel('Predicted Sale Price')
plt.title('ElasticNet: Sale Price vs Predicted Sale Price')

# Residuals

plt.subplot(1, 2, 2)
sns.residplot(x=y_log_scale, y=y_pred)
plt.xlabel('Sale Price')
plt.ylabel('Residual')
plt.title('Residual Plot')

MSE = np.square(np.subtract(y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_scale, y_pred)
print("R_sq:",r2)

```


```

Out[76]: ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,
max_iter=1000, normalize=False, positive=False, precompute=False,
random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[76]: ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,
max_iter=1000, normalize=False, positive=False, precompute=False,
random_state=0, selection='cyclic', tol=0.0001, warm_start=False)
[ 0.         0.         0.17296878 -0.         0.         0.
  0.         0.         0.00587394  0.         0.        -0.
  0.06700516  0.        -0.         0.         0.         0.
 -0.         0.         0.         0.02809537  0.00115474  0.
  0.        -0.         0.         0.         0.        -0.        ]
[-2.46470368e-15]
Out[76]: <AxesSubplot:>
Out[76]: <matplotlib.collections.PathCollection at 0x7fee824f0f10>
Out[76]: Text(0.5, 0, 'Sale Price (Standard Scaled)')
Out[76]: Text(0, 0.5, 'Predicted Sale Price')
Out[76]: Text(0.5, 1.0, 'ElasticNet: Sale Price vs Predicted Sale Price')
Out[76]: <AxesSubplot:>
Out[76]: <AxesSubplot:>
Out[76]: Text(0.5, 0, 'Sale Price')
Out[76]: Text(0, 0.5, 'Residual')
Out[76]: Text(0.5, 1.0, 'Residual Plot')

```

MSE: 1.058250120091334

R_sq: 0.36078782270420806



```
In [77]: x = housing_training_data[['OverallQual', 'TotalBsmtSF', 'GrLivArea', 'GarageCars', 'GarageArea']]
y_log = np.log(housing_training_data['SalePrice'])

#polynomial_features= PolynomialFeatures(degree=2)
#xp = polynomial_features.fit_transform(x)
#xp.shape

#add constant to predictor variables
#x = sm.add_constant(xp)
x = sm.add_constant(x)

#fit polynomial regression model
model = sm.OLS(y_log, x).fit()

#view model summary
print(model.summary())
```

OLS Regression Results

Dep. Variable:	SalePrice	R-squared:	0.825			
Model:	OLS	Adj. R-squared:	0.825			
Method:	Least Squares	F-statistic:	1370.			
Date:	Sun, 16 Apr 2023	Prob (F-statistic):	0.00			
Time:	16:12:42	Log-Likelihood:	555.36			
No. Observations:	1455	AIC:	-1099.			
Df Residuals:	1449	BIC:	-1067.			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	10.5342	0.020	524.743	0.000	10.495	10.574
OverallQual	0.1194	0.005	25.963	0.000	0.110	0.128
TotalBsmstSF	0.0002	1.3e-05	15.556	0.000	0.000	0.000
GrLivArea	0.0002	1.11e-05	21.632	0.000	0.000	0.000
GarageCars	0.0689	0.013	5.246	0.000	0.043	0.095
GarageArea	0.0001	4.51e-05	3.118	0.002	5.22e-05	0.000
=====						
Omnibus:	321.730	Durbin-Watson:	1.974			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1037.251			
Skew:	-1.085	Prob(JB):	5.80e-226			
Kurtosis:	6.522	Cond. No.	9.24e+03			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.24e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Polynomial regression using predictors from PCA for Kaggle

```
In [78]: from sklearn.preprocessing import PolynomialFeatures
ypolyscaler = StandardScaler()
poly_y_scale = ypolyscaler.fit_transform(np.array(housing_training_data['SalePrice'])).
#define our polynomial model, with whatever degree we want
degree=2

# PolynomialFeatures will create a new matrix consisting of all polynomial combination
# of the features with a degree less than or equal to the degree we just gave the model
poly_model = PolynomialFeatures(degree=degree)

# transform out polynomial features
poly_x_values = poly_model.fit_transform(x_pca_scale)

poly_regression_model = LinearRegression()
poly_regression_model.fit(poly_x_values, poly_y_scale)
y_pred = poly_regression_model.predict(poly_x_values)

plt.subplot(1,2,1)

plt.scatter(poly_y_scale, y_pred)
plt.xlabel('Sale Price (Standard Scaled)')
plt.ylabel('Predicted Sale Price')
plt.title('Sale Price vs Predicted Sale Price')
```

```
# Residuals
plt.subplot(1,2,2)
plt.title('Residuals')
sns.residplot(x=poly_y_scale, y=y_pred)

# Mean Squared Error

MSE = np.square(np.subtract(poly_y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_scale, y_pred)
print("R_sq:",r2)
```

Out[78]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

Out[78]: <AxesSubplot:>

Out[78]: <matplotlib.collections.PathCollection at 0x7fee826cf9d0>

Out[78]: Text(0.5, 0, 'Sale Price (Standard Scaled)')

Out[78]: Text(0, 0.5, 'Predicted Sale Price')

Out[78]: Text(0.5, 1.0, 'Sale Price vs Predicted Sale Price')


Out[78]: <AxesSubplot:>

Out[78]: Text(0.5, 1.0, 'Residuals')

Out[78]: <AxesSubplot:title={'center': 'Residuals'}>

MSE: 0.08764742224548526

R_sq: 0.9123525777545147



Ridge regression using predictors from PCA with Cross Validation to find the best alpha

Principal components have no collinearity by definition. However, we were interested in applying a Ridge regularization model to our PCA and seeing how it would perform, including what value of alpha it would select. There are many methods aimed at making our PCA more robust/reduce overfitting inherent in their calculation. See below for comments about the loading within our components.

```

In [79]: from sklearn.model_selection import GridSearchCV
yscaler = StandardScaler()
y_log_scale = yscaler.fit_transform(np.array(np.log(housing_training_data['SalePrice']

pca = PCA(n_components=8)

## independent variables ###
xscaler = StandardScaler()
x_raw = housing_training_data.select_dtypes(exclude=['object']).drop(columns = ['SaleP
x_scale = xscaler.fit_transform(x_raw)

x_pca_raw = pca.fit_transform(x_raw)

x_pca_scale = pca.fit_transform(x_scale) # PCA is affected by scale. We use the scaled

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x_pca_scale, y_log_scale, test_siz

# Set the alpha values to test
alpha_values = np.logspace(-4, 4, num=50)

# Create the Ridge Regression model
ridge = Ridge()

# Set up the grid search with cross-validation
param_grid = {'alpha': alpha_values}
grid_search = GridSearchCV(ridge, param_grid, cv=5, scoring='neg_mean_squared_error',

# Fit the grid search to the training data
grid_search.fit(X_train, y_train)

# Get the best alpha value
best_alpha = grid_search.best_params_['alpha']
print("Best alpha value:", best_alpha)

# Create and fit the Ridge Regression model with the best alpha value
ridge_best = Ridge(alpha=best_alpha)

ridgemodel = ridge_best.fit(X_train, y_train)
# View the coefficients of the ridge model
coefficients = ridgemodel.coef_
print("Coefficients:", coefficients)

# Predict
y_pred = ridgemodel.predict(X_test)

# plot predictors against log scaled
plt.scatter(y_test, y_pred)
plt.xlabel('Sale Price (Standard Scaled)')
plt.ylabel('Predicted Sale Price')
plt.title('Sale Price vs Predicted Sale Price')
# Calculate the mean squared error (MSE) of the predictions
MSE = np.square(np.subtract(y_test,y_pred)).mean()
print("MSE:",MSE)

# Calculate the R^2 of the predictions
r2 = r2_score(y_test, y_pred)
print("R_sq:",r2)

```

```

Out[79]: GridSearchCV(cv=5, error_score=nan,
                    estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                   max_iter=None, normalize=False, random_state=None,
                                   solver='auto', tol=0.001),
                    iid='deprecated', n_jobs=-1,
                    param_grid={'alpha': array([1.00000000e-04, 1.45634848e-04, 2.12095089e-
04, 3.08884360e-04,
                    4.49843267e-04, 6.55128557e-04, 9.54095476e-04, 1.38949549e-03,
                    2.02358965e-03, 1.67683294e+01, 2.44205309e+01, 3.55648031e+01, 5.17947468e+01,
                    7.54312006e+01, 1.09854114e+02, 1.59985872e+02, 2.32995181e+02,
                    3.39322177e+02, 4.94171336e+02, 7.19685673e+02, 1.04811313e+03,
                    1.52641797e+03, 2.22299648e+03, 3.23745754e+03, 4.71486636e+03,
                    6.86648845e+03, 1.00000000e+04])},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring='neg_mean_squared_error', verbose=0)
Best alpha value: 24.420530945486497
Coefficients: [[-0.28015625  0.04133806  0.05587674 -0.04646577  0.05143315 -0.078062
32
               0.01581126 -0.01322861]]
Out[79]: <matplotlib.collections.PathCollection at 0x7fee87819e50>
Out[79]: Text(0.5, 0, 'Sale Price (Standard Scaled)')
Out[79]: Text(0, 0.5, 'Predicted Sale Price')
Out[79]: Text(0.5, 1.0, 'Sale Price vs Predicted Sale Price')
MSE: 0.12096607983798188
R_sq: 0.8936548969444329

```



```

In [80]: x_raw.columns

```

```

Out[80]: Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
               'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'ExterQual',
               'BsmtQual', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
               '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath',
               'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
               'KitchenQual', 'TotRmsAbvGrd', 'Fireplaces', 'FireplaceQu',
               'GarageYrBlt', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
               'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
               'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'Street_Pave',
               'CentralAir_Y', 'PavedDrive_Y', 'HeatingEx', 'YrSinceRemod', 'TotalSF'],
              dtype='object')

```


Interpret the coefficients

Coef0: (0.29916159) has a positive impact on Sale Price. Given that As this feature increases, the predicted house price also increases. This coefficient has the largest positive effect among all features. In the PCA loadings for PC1, TotalSF and OverQual accounting for over 0.50 of the component, indicating both these features have a larger positive impact on Sale Price.

Coef1: (0.01065569) has a small positive impact on Sale Price. As this feature increases, the predicted house price marginally increases.

Coef2: (0.05379505) has a positive impact on Sale Price. As this feature increases, the predicted house price increases.

Coef3: (0.08326186) has a positive impact on Sale Price. As this feature increases, the predicted house price increases.

Coef4: (-0.09188438) has a negative impact on Sale Price. As this feature increases, the predicted house price decreases. This coefficient has the largest negative effect among all features.

Coef5: (0.03202796) has a positive impact on Sale Price. As this feature increases, the predicted house price increases.


Coef6: (0.01423303) has a small positive impact on Sale Price. As this feature increases, the predicted house price marginally increases.

Coef7: (-0.01638527) has a small negative impact on Sale Price. As this feature increases, the predicted house price marginally decreases.

In [81]: `# Residuals`

```
sns.residplot(x=y_test, y=y_pred)
```

Out[81]: `<AxesSubplot:>`



Lasso Regression

We will now try to fit a Lasso Regression to the housing sales dataframe.

```
In [82]: # Import libraries relevant to Lasso
from sklearn.linear_model import LassoCV
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
import pandas as pd

# Create dataframe that can be used for Lasso. Only keep variables that aren't causing
lasso_sandbox = housing_training_data[ ['LotFrontage', 'LotArea', 'OverallQual', 'Over
    'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '
    'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
    'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF
    'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', '

# Perform a Log Transformation on the outcome variable SalePrice and drop the original
lasso_sandbox_log_saleprice = np.log(housing_training_data['SalePrice'])
lasso_sandbox_x = lasso_sandbox.drop(columns=['SalePrice'])

# Split the dataset into training and testing dataframes
lasso_X_train, lasso_X_validation, lasso_y_train, lasso_y_validation = train_test_split(
    lasso_sandbox, lasso_y_train, random_state=42)

# Standardize the numeric predictors - which can help strengthen the model fit
numerical_predictors = ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'MasVr
    'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '
    'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
    'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF
    'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']

from sklearn.preprocessing import StandardScaler
lasso_scaler = StandardScaler().fit(lasso_X_train[numerical_predictors])
lasso_X_train[numerical_predictors] = lasso_scaler.transform(lasso_X_train[numerical_p
lasso_X_validation[numerical_predictors] = lasso_scaler.transform(lasso_X_validation[

# Let's visualize the Lasso coefficients as a function of the tuning parameter, alpha
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Lasso

alphas = np.linspace(0.01, 500, 100)
lasso = Lasso(max_iter=10000)
coefs = []

for a in alphas:
    lasso.set_params(alpha=a)
    lasso.fit(lasso_X_train, lasso_y_train)
    coefs.append(lasso.coef_)

ax = plt.gca()

ax.plot(alphas, coefs)
```

```

ax.set_xscale('log')
plt.axis('tight')
plt.xlabel('alpha')
plt.ylabel('Standardized Coefficients')
plt.title('Lasso coefficients as a function of alpha');

# Fit a Lasso Regression Model with ten-fold cross-validation
lasso_model = LassoCV(cv=5, random_state=1, max_iter = 10000)
lasso_model.fit(lasso_X_train, lasso_y_train)

lasso_model.coef_

```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/pandas/core/frame.py:3678: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[col] = igetitem(value, i)
```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/pandas/core/frame.py:3678: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self[col] = igetitem(value, i)
```

```

Out[82]: Lasso(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=5.06040404040404, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=5.06040404040404, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=10.11080808080808, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=10.11080808080808, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=15.16121212121212, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=15.16121212121212, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=20.21161616161616, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=20.21161616161616, copy_X=True, fit_intercept=True, max_iter=10000,
              normalize=False, positive=False, precompute=False, random_state=None,
              selection='cyclic', tol=0.0001, warm_start=False)

```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]


[illegible]

```
Out[82]: Lasso(alpha=479.7983838383838, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=479.7983838383838, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=484.8487878787879, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=484.8487878787879, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=489.89919191919194, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=489.89919191919194, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=494.94959595959597, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=494.94959595959597, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=500.0, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: Lasso(alpha=500.0, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[82]: [<matplotlib.lines.Line2D at 0x7fee87dd4210>,
          <matplotlib.lines.Line2D at 0x7fee87dd4450>,
          <matplotlib.lines.Line2D at 0x7fee87dd4590>,
          <matplotlib.lines.Line2D at 0x7fee87dd4690>,
          <matplotlib.lines.Line2D at 0x7fee87dd4790>,
          <matplotlib.lines.Line2D at 0x7fee87dd47d0>,
          <matplotlib.lines.Line2D at 0x7fee87dd4a10>,
          <matplotlib.lines.Line2D at 0x7fee87dd4b10>,
          <matplotlib.lines.Line2D at 0x7fee87dd4c10>,
          <matplotlib.lines.Line2D at 0x7fee87dd48d0>,
          <matplotlib.lines.Line2D at 0x7fee87d62b90>,
          <matplotlib.lines.Line2D at 0x7fee87d8d110>,
          <matplotlib.lines.Line2D at 0x7fee87dd4e50>,
          <matplotlib.lines.Line2D at 0x7fee87dd4f50>,
          <matplotlib.lines.Line2D at 0x7fee87dd4f90>,
          <matplotlib.lines.Line2D at 0x7fee87de1190>,
          <matplotlib.lines.Line2D at 0x7fee87de1290>,
          <matplotlib.lines.Line2D at 0x7fee87de1390>,
          <matplotlib.lines.Line2D at 0x7fee87de1490>,
          <matplotlib.lines.Line2D at 0x7fee87de1590>,
          <matplotlib.lines.Line2D at 0x7fee87de1690>,
          <matplotlib.lines.Line2D at 0x7fee87de1790>,
          <matplotlib.lines.Line2D at 0x7fee87de1890>,
          <matplotlib.lines.Line2D at 0x7fee87de1990>,
          <matplotlib.lines.Line2D at 0x7fee87de1a90>,
          <matplotlib.lines.Line2D at 0x7fee87de1b90>,
          <matplotlib.lines.Line2D at 0x7fee87de1c90>,
          <matplotlib.lines.Line2D at 0x7fee87de1d90>,
          <matplotlib.lines.Line2D at 0x7fee87de1e90>,
          <matplotlib.lines.Line2D at 0x7fee87de1f90>]
```

```

Out[82]: (0.005821722491960459,
          858.8523425678187,
          -0.028034197659991166,
          0.1415390558301679)
Out[82]: Text(0.5, 0, 'alpha')
Out[82]: Text(0, 0.5, 'Standardized Coefficients')
Out[82]: Text(0.5, 1.0, 'Lasso coefficients as a function of alpha')
Out[82]: LassoCV(alphas=None, copy_X=True, cv=5, eps=0.001, fit_intercept=True,
                  max_iter=10000, n_alphas=100, n_jobs=None, normalize=False,
                  positive=False, precompute='auto', random_state=1, selection='cyclic',
                  tol=0.0001, verbose=False)
Out[82]: array([ 0.0144498 ,  0.01915226,  0.13032816,  0.03233272,  0.          ,
                  0.0253451 ,  0.          , -0.          ,  0.05136664,  0.02051305,
                  0.          , -0.00935478,  0.07656874,  0.02401095,  0.          ,
                  0.04353805,  0.02686431, -0.          , -0.02599228,  0.          ,
                  0.01277892,  0.05172721,  0.0221585 ,  0.01390011,  0.01036037,
                  -0.00643844, -0.          ,  0.00428796,  0.          ,  0.          ])

```



Let's plot the mean squared error as a function of our tuning parameter, alpha, from our cross-validation

```

In [83]: plt.semilogx(lasso_model.alphas_, lasso_model.mse_path_, ":")
plt.plot(
    lasso_model.alphas_,
    lasso_model.mse_path_.mean(axis=-1),
    "k",
    label="Average across the folds",
    linewidth=2,
)
plt.axvline(
    lasso_model.alpha_, linestyle="--", color="k", label="alpha: CV estimate"
)


plt.legend()
plt.xlabel("alphas")
plt.ylabel("Mean square error")
plt.title("Mean square error on each fold")
plt.axis("tight")

```

```

Out[83]: [

```



Let's fit a Lasso Regression using the newly identified best value for the tuning parameter, alpha.

```

In [84]: alphas = np.linspace(0.01,500,100)

best_alpha
alphas[0:10]

```

```

Out[84]: 24.420530945486497

```

```

Out[84]: array([1.00000000e-02, 5.06040404e+00, 1.01108081e+01, 1.51612121e+01,
2.02116162e+01, 2.52620202e+01, 3.03124242e+01, 3.53628283e+01,
4.04132323e+01, 4.54636364e+01])

```

```

In [85]: ### plot training vs testing MSE for increasing lambda

```

```

from sklearn.metrics import mean_squared_error

lasso = Lasso(max_iter=10000)
alphas = np.linspace(0.00005,0.015,100)
#alphas = np.linspace(0.00001,0.00003,1000)

scaler = StandardScaler()
lasso_x = scaler.fit_transform(lasso_sandbox_x)
lasso_y = scaler.fit_transform(np.array(lasso_sandbox['SalePrice']).reshape(-1,1))

```

```

# Split the dataset into training and testing dataframes
lasso_X_train, lasso_X_validation, lasso_y_train, lasso_y_validation = train_test_split(
    lasso_X, lasso_y, test_size=0.2, random_state=42)

training_mse = []
testing_mse = []

#best_alpha = lasso_model.alpha_

for a in alphas:
    lasso.set_params(alpha=a)
    lasso.fit(lasso_X_train, lasso_y_train)
    training_mse.append(mean_squared_error(lasso_y_train, lasso.predict(lasso_X_train)))
    testing_mse.append(mean_squared_error(lasso_y_validation, lasso.predict(lasso_X_validation)))
#coefs.append(Lasso.coef_)

```

```

Out[85]: Lasso(alpha=5e-05, copy_X=True, fit_intercept=True, max_iter=10000,
           normalize=False, positive=False, precompute=False, random_state=None,
           selection='cyclic', tol=0.0001, warm_start=False)

```

```

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn
n/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not conv
erge. You might want to increase the number of iterations. Duality gap: 0.33398039476
475105, tolerance: 0.10374553746774125
    positive)

```

```

Out[85]: Lasso(alpha=5e-05, copy_X=True, fit_intercept=True, max_iter=10000,
           normalize=False, positive=False, precompute=False, random_state=None,
           selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.000201010101010101, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.000201010101010101, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00035202020202020203, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00035202020202020203, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00050303030303030303, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00050303030303030303, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00065404040404040404, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00065404040404040404, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00080505050505050505, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```

```

Out[85]: Lasso(alpha=0.00080505050505050505, copy_X=True, fit_intercept=True,
           max_iter=10000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)

```


[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```

Out[85]: Lasso(alpha=0.014546969696969695, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.014546969696969695, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.014697979797979797, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.014697979797979797, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.014848989898989898, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.014848989898989898, copy_X=True, fit_intercept=True,
          max_iter=10000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.015, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)
Out[85]: Lasso(alpha=0.015, copy_X=True, fit_intercept=True, max_iter=10000,
          normalize=False, positive=False, precompute=False, random_state=None,
          selection='cyclic', tol=0.0001, warm_start=False)

```

```

In [86]: diff = np.array(testing_mse) - np.array(training_mse)
          #diff

```


```

In [87]: plt.plot(alphas, training_mse, label='Training Error')
          plt.plot(alphas, testing_mse, label='Testing Error')
          #plt.axvline(best_alpha, linestyle="--", color="k", label="Optimal Alpha")
          plt.legend()

          ax.set_xscale('log')
          plt.axis('tight')
          plt.xlabel('Alpha')
          plt.ylabel('Mean Squared Error')
          plt.title('Training MSE vs Testing MSE');

          # Fit a Lasso Regression Model with ten-fold cross-validation
          #lasso_model = LassoCV(cv=5, random_state=1, max_iter = 10000)
          #lasso_model.fit(lasso_X_train, lasso_y_train)

```



```

In [88]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score, get_scorer
from sklearn.linear_model import Lasso, Ridge, LassoCV, LinearRegression, ElasticNet
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.model_selection import KFold, RepeatedKFold, GridSearchCV, cross_validate

def regmodel_param_test(
    alphas_to_try, X, y, cv, scoring = 'r2',
    model_name = 'LASSO', X_test = None, y_test = None,
    draw_plot = False, filename = None):

    validation_scores = []
    train_scores = []
    results_list = []
    if X_test is not None:
        test_scores = []
        scorer = get_scorer(scoring)
    else:
        test_scores = None

    for curr_alpha in alphas_to_try:

        if model_name == 'LASSO':
            regmodel = Lasso(alpha = curr_alpha)
        elif model_name == 'Ridge':
            regmodel = Ridge(alpha = curr_alpha)
        elif model_name == 'ElasticNet':
            regmodel = ElasticNet(alpha = curr_alpha)
        else:
            return None

        results = cross_validate(
            regmodel, X, y, scoring=scoring, cv=cv,
            return_train_score = True)

        validation_scores.append(np.mean(results['test_score']))
        train_scores.append(np.mean(results['train_score']))
        results_list.append(results)

        if X_test is not None:
            regmodel.fit(X,y)
            y_pred = regmodel.predict(X_test)
            test_scores.append(scorer(regmodel, X_test, y_test))

    chosen_alpha_id = np.argmax(validation_scores)
    chosen_alpha = alphas_to_try[chosen_alpha_id]
    max_validation_score = np.max(validation_scores)
    if X_test is not None:
        test_score_at_chosen_alpha = test_scores[chosen_alpha_id]
    else:
        test_score_at_chosen_alpha = None

    if draw_plot:
        regmodel_param_plot(
            validation_scores, train_scores, alphas_to_try, chosen_alpha,
            scoring, model_name, test_scores, filename)

```

```

    return chosen_alpha, max_validation_score, test_score_at_chosen_alpha

def regmodel_param_plot(
    validation_score, train_score, alphas_to_try, chosen_alpha,
    scoring, model_name, test_score = None, filename = None):

    plt.figure(figsize = (8,8))
    sns.lineplot(y = validation_score, x = alphas_to_try,
                 label = 'validation_data')
    sns.lineplot(y = train_score, x = alphas_to_try,
                 label = 'training_data')
    plt.axvline(x=chosen_alpha, linestyle='--')
    if test_score is not None:
        sns.lineplot(y = test_score, x = alphas_to_try,
                     label = 'test_data')
    plt.xlabel('alpha_parameter')
    plt.ylabel(scoring)
    plt.title(model_name + ' Regularisation')
    plt.legend()
    if filename is not None:
        plt.savefig(str(filename) + ".png")
    plt.show()

```

```

In [89]: cv = KFold(n_splits=5, shuffle=True)

lasso_alphas = np.linspace(0, 0.02, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, lasso_X_train, lasso_y_train,
        cv, scoring = 'r2', model_name = 'LASSO',
        X_test = lasso_X_validation, y_test = lasso_y_validation,
        draw_plot = True, filename = 'lasso_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)

```


```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 59.31511840624592, tolerance: 0.08146137472628749
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 62.41338009321754, tolerance: 0.08383375419276164
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 52.001325039859324, tolerance: 0.07646279640865453
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 65.00951846597536, tolerance: 0.08599641085489657
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 64.78763983418031, tolerance: 0.08721964620031875
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:44: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 77.05558781228252, tolerance: 0.10374553746774125
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.34496764577092165, tolerance: 0.08126724283571046
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.12904315448244574, tolerance: 0.07752593667652338
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.40067582136090607, tolerance: 0.08818182766329767
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.211249324772524, tolerance: 0.08300929344027744
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.4239314089955144, tolerance: 0.08483106610569778
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.36810356882534734, tolerance: 0.10374553746774125
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.2361188921867381, tolerance: 0.08262500137838019
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.40371434196553935, tolerance: 0.07940128899327321
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.20887390371
```

```

1023, tolerance: 0.08018441167777415
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.4938553561466392, tolerance: 0.083430130263864
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.2495814785760473, tolerance: 0.0893172914193282
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.39077165557705484, tolerance: 0.10374553746774125
positive)

```



Chosen alpha: 0.00600

Validation score: 0.84299

Test score at chosen alpha: 0.86357

```

In [90]: ### BASED ON LASSOCV:

# Show best value of tuning parameter (alpha) chosen by cross validation
print(f"The best value of the tuning parameter, alpha, chosen by cross-validation is {

```

```
# Set best tuning parameter (alpha) and use it to fit our final Lasso regression model
lasso_best = Lasso(alpha=lasso_model.alpha_)
lasso_best.fit(lasso_X_train, lasso_y_train)

# Show best Lasso model coefficients and names
best_lasso_coeffs_data = (list(zip(lasso_X_train, lasso_best.coef_)))

best_lasso_coeffs_df = pd.DataFrame(best_lasso_coeffs_data, columns = ['Predictor', 'E

with pd.option_context("display.max_rows", 300):
    best_lasso_coeffs_df.style.background_gradient(cmap = 'Greens')
```

The best value of the tuning parameter, alpha, chosen by cross-validation is 0.0035778259276540722.

```
Out[90]: Lasso(alpha=0.0035778259276540722, copy_X=True, fit_intercept=True,
          max_iter=1000, normalize=False, positive=False, precompute=False,
          random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
```


Out[90]:

	Predictor	Best Lasso Coefficient
0	[-0.60134894 -0.12192403 -0.7943627 -0.51768686 -0.57401865 0.58069606 -0.28918978 -0.59461551 -0.14495197 -0.45356901 -0.79680412 -0.12045114 -1.04155359 -0.81857737 3.96564919 -1.02643215 -0.75919202 0.16691984 -0.21183287 -0.93449284 -0.95104897 0.31592367 0.94577272 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.041777
1	[-0.93621711 -0.45101191 -0.06376995 0.38008952 -0.57401865 -0.09178658 -0.28918978 0.31764015 0.13181675 -0.14422762 -0.79680412 -0.12045114 -0.81315996 -0.81857737 -0.23990443 -1.02643215 -0.75919202 -1.05968807 -0.21183287 -0.93449284 0.61077524 0.31592367 -0.1486834 0.577744 1.12931613 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.067138
2	[-0.02728921 -0.13358662 -0.06376995 2.17564227 0.84977655 -0.08480579 0.10705153 -0.52199815 -0.61031529 -0.97370056 1.13164459 -0.12045114 0.25604563 -0.81857737 3.96564919 -1.02643215 1.23112219 1.39352774 -0.21183287 0.30836555 -0.95104897 0.31592367 0.05888586 -0.74905884 -0.0948976 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.321800
3	[0.4510939 0.01924402 1.39741555 -0.51768686 1.32437495 1.98150764 -0.28918978 -0.33818544 1.60628371 1.53935609 -0.79680412 -0.12045114 0.42986733 1.114913 -0.23990443 0.80378694 -0.75919202 0.16691984 -0.21183287 0.30836555 0.61077524 1.65606634 1.67226514 0.91344111 -0.03368691 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.064117
4	[0.11622572 -0.1832794 0.6668228 -0.51768686 -0.57401865 -1.0155776 -0.28918978 1.82898909 0.79067325 0.59216916 -0.79680412 -0.12045114 -0.26946184 -0.81857737 -0.23990443 0.80378694 -0.75919202 0.16691984 -0.21183287 -0.31306364 -0.95104897 0.31592367 0.27117261 -0.74905884 1.43536957 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.073773
5	[0.92947701 0.20422285 1.39741555 -0.51768686 -0.57401865 2.37243181 -0.28918978 -0.90323932 1.40789198 1.28202785 -0.79680412 -0.12045114 0.23987617 1.114913 -0.23990443 0.80378694 -0.75919202 -1.05968807 -0.21183287 -0.93449284 0.61077524 1.65606634 1.3467588 -0.74905884 1.0375001 2.79957949 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.092032
6	[-0.93621711 -0.4729173 -0.7943627 2.17564227 -0.57401865 -1.0155776 -0.28918978 -0.85331488 -2.1043765 -0.7382726 -0.79680412 -0.12045114 -1.25175658 -0.81857737 -0.23990443 -1.02643215 -0.75919202 -1.05968807 -0.21183287 -1.55592204 -0.95104897 -1.024219 -0.93650311 0.40190266 -0.40095103 0.95699807 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.002737
7	[-0.21864245 -0.26806136 -1.52495546 -2.31323961 -0.57401865 0.43177257 -0.28918978 -0.14983414 0.17835308 0.47171764 0.71674998 -0.12045114 0.96143835 1.114913 -0.23990443 0.80378694 -0.75919202 1.39352774 -0.21183287 0.30836555 2.17259946 0.31592367 0.49289432 2.69583288 -0.03368691 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.000000
8	[-0.12296583 0.19337157 -0.06376995 2.17564227 -0.57401865 -1.0155776 -0.28918978 0.26771571 -0.89443185 -1.29125456 0.38993355 -0.12045114 -0.62518997 -0.81857737 -0.23990443 -1.02643215 -0.75919202 0.16691984 -0.21183287 0.30836555 -0.95104897 0.31592367 0.26645513 -0.74905884 -0.00308157 -0.35913151 -0.11654172 -0.27070619 -0.05813532 1.5224731]	0.139404
9	[-0.21864245 -0.20254802 -0.7943627 0.38008952 -0.57401865 0.78779279 -0.28918978 -0.64680925 0.01670056 -0.27289174 -0.79680412 -0.12045114 -0.90815554 1.114913 -0.23990443 -1.02643215 -0.75919202 0.16691984 -0.21183287 -0.31306364 0.61077524 -1.024219 -0.79026113 -0.74905884 0.59372263 2.66796653 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.041091

	Predictor	Best Lasso Coefficient
10	[-0.31431907 -0.26086098 0.6668228 -0.51768686 -0.35931937 -1.0155776 -0.28918978 1.52263458 0.46002036 0.26640256 -0.79680412 -0.12045114 -0.50998257 -0.81857737 -0.23990443 0.80378694 -0.75919202 0.16691984 -0.21183287 -0.31306364 0.61077524 0.31592367 -0.32794777 1.00935457 -0.38564836 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.000000
11	[-0.79270218 -0.68588634 0.6668228 -0.51768686 0.16047887 -1.0155776 -0.28918978 1.83125838 0.79312253 1.09313797 -0.79680412 -0.12045114 0.10041457 -0.81857737 -0.23990443 0.80378694 -0.75919202 -1.05968807 -0.21183287 0.30836555 0.61077524 1.65606634 0.88916292 0.39390987 -0.40095103 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.025678
12	[0.25974066 -0.06310401 -0.7943627 -0.51768686 -0.57401865 -1.0155776 -0.28918978 -1.2844805 -2.56973983 1.39152923 -0.79680412 -0.12045114 0.32072348 -0.81857737 -0.23990443 0.80378694 -0.75919202 1.39352774 4.32076744 0.92979475 -0.95104897 -2.36436167 -2.22437604 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.253948
13	[-0.98405542 -0.4692664 -2.25554821 2.17564227 -0.57401865 -0.41988364 -0.28918978 0.77149869 0.27632431 0.02002445 -0.79680412 -0.12045114 -0.691889 1.114913 -0.23990443 -1.02643215 -0.75919202 0.16691984 -0.21183287 -0.31306364 -0.95104897 -1.024219 -1.18652972 0.38591708 0.79265736 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.022657
14	[-0.93621711 -0.45101191 -0.7943627 2.17564227 -0.57401865 -1.0155776 -0.28918978 0.72157425 -0.40457572 -0.74374767 0.2786769 -0.12045114 -0.31797022 1.114913 -0.23990443 -1.02643215 -0.75919202 0.16691984 -0.21183287 -0.93449284 0.61077524 -1.024219 -1.20539966 -0.74905884 -0.70700446 3.0628054 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.016310
15	[0.30757897 -0.28570737 -0.7943627 3.07341864 -0.57401865 -1.0155776 -0.28918978 -0.46753512 -1.68799879 -0.33859256 1.01111655 -0.12045114 0.6198585 -0.81857737 -0.23990443 0.80378694 -0.75919202 1.39352774 4.32076744 0.92979475 0.61077524 0.31592367 0.94577272 2.0004603 -0.70700446 0.29893328 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.061960
16	[-0.45783401 0.03577448 0.6668228 2.17564227 -0.57401865 -1.0155776 -0.28918978 0.34487166 -0.81115631 1.15062619 1.47005026 -0.12045114 2.11957598 -0.81857737 -0.23990443 -1.02643215 1.23112219 0.16691984 -0.21183287 0.92979475 -0.95104897 0.31592367 1.09673218 -0.74905884 0.02752377 1.99345012 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.034481
17	[0.49893221 0.25726228 -0.7943627 0.38008952 2.01932261 1.969873 0.62092699 -1.2844805 0.93273153 1.34499114 -0.79680412 -0.12045114 0.28636337 1.114913 -0.23990443 0.80378694 -0.75919202 0.16691984 -0.21183287 0.30836555 0.61077524 0.31592367 0.70518107 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.120404
18	[0.4510939 -0.09190553 0.6668228 -0.51768686 0.94017625 -1.0155776 -0.28918978 2.12399714 1.10907974 0.94804864 -0.79680412 -0.12045114 -0.0067081 -0.81857737 -0.23990443 0.80378694 -0.75919202 0.16691984 -0.21183287 0.30836555 -0.95104897 0.31592367 0.81368319 -0.74905884 1.0375001 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.076264
19	[0.92947701 0.02360482 -0.7943627 -2.31323961 -0.32541949 0.74358113 -0.28918978 1.84714343 2.66192368 2.68364551 -0.79680412 -0.12045114 1.27472166 -0.81857737 -0.23990443 0.80378694 -0.75919202 1.39352774 -0.21183287 0.30836555 -0.95104897 0.31592367 0.26645513 -0.74905884 -0.24792432 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.106356

	Predictor	Best Lasso Coefficient
20	[0.49893221 0.63330477 0.6668228 -0.51768686 0.46557785 0.82269673 2.61451608 -0.98266457 0.83965886 0.64691985 2.42963892 -0.12045114 2.58444797 1.114913 -0.23990443 2.63400602 1.23112219 2.62013565 -0.21183287 3.41551155 2.17259946 0.31592367 0.43628453 -0.74905884 1.00689476 -0.35913151 -0.11654172 7.6107848 14.60962252 3.93794201]	-0.000000
21	[0.49893221 -0.1928123 -1.52495546 -0.51768686 -0.57401865 -1.0155776 -0.28918978 1.25939662 0.1759038 -0.094952 -0.79680412 -0.12045114 -0.77677867 -0.81857737 -0.23990443 0.80378694 -0.75919202 -1.05968807 -0.21183287 -0.93449284 -0.95104897 0.31592367 -0.1486834 0.3059892 0.27236652 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.065079
22	[-0.64918725 1.01249108 2.12800831 -0.51768686 4.34146478 2.26539305 -0.28918978 -0.29053029 1.95653085 1.89523557 1.24290125 -0.12045114 2.47126175 1.114913 -0.23990443 0.80378694 1.23112219 1.39352774 -0.21183287 2.17265315 2.17259946 1.65606634 1.10616715 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.064081
23	[1.12083025 0.07228346 0.6668228 -0.51768686 -0.57401865 1.46260236 -0.28918978 1.08693038 2.59824238 2.61246961 -0.79680412 -0.12045114 1.22217091 1.114913 -0.23990443 0.80378694 1.23112219 0.16691984 -0.21183287 0.92979475 2.17259946 0.31592367 0.23815023 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.030624
24	[-1.22324698 -0.41338738 0.6668228 -0.51768686 0.14917891 0.69006175 -0.28918978 0.13382744 0.75638332 0.55384368 -0.79680412 -0.12045114 -0.2977584 1.114913 -0.23990443 0.80378694 -0.75919202 -1.05968807 -0.21183287 -0.31306364 0.61077524 0.31592367 0.05888586 0.7855565 -0.17141096 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.028564
25	[0.0205491 -0.2551818 -0.7943627 -0.51768686 -0.57401865 1.65573751 -0.28918978 -1.2844805 0.24203438 -0.02103856 -0.79680412 -0.12045114 -0.72220674 1.114913 -0.23990443 -1.02643215 -0.75919202 0.16691984 -0.21183287 -0.31306364 -0.95104897 -1.024219 0.94577272 -0.74905884 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.022844
26	[-0.55351063 0.64121505 0.6668228 -0.51768686 -0.40451922 -1.0155776 -0.28918978 1.42732428 0.35715057 0.10762556 0.69588936 -0.12045114 0.67443043 -0.81857737 -0.23990443 0.80378694 1.23112219 1.39352774 -0.21183287 0.30836555 -0.95104897 0.31592367 0.06832083 -0.74905884 0.53251194 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.013004
27	[-0.02728921 -0.02040878 -1.52495546 3.07341864 -0.57401865 0.49227274 -0.28918978 -1.2844805 -0.98260595 -0.53569505 -0.79680412 -0.12045114 -1.10218906 1.114913 3.96564919 -2.85665123 -0.75919202 -3.51290387 -0.21183287 -2.17735124 -0.95104897 -1.024219 -0.50721213 -0.04569348 -0.70700446 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	0.012988
28	[-0.26648076 -0.35172777 -0.06376995 -0.51768686 -0.57401865 1.44165999 -0.28918978 -1.2844805 0.01670056 -0.27289174 -0.79680412 -0.12045114 -0.90815554 3.04840337 -0.23990443 -2.85665123 -0.75919202 -3.51290387 4.32076744 -1.55592204 -0.95104897 0.31592367 0.49289432 1.36103725 0.14994515 -0.35913151 -0.11654172 -0.27070619 -0.05813532 1.11989495]	0.000040
29	[-0.07512752 -0.16238815 -0.06376995 -0.51768686 -0.57401865 -1.0155776 -0.28918978 0.41748903 -0.73277933 -1.11057729 0.94158114 -0.12045114 -0.01075047 -0.81857737 -0.23990443 0.80378694 1.23112219 0.16691984 -0.21183287 -0.31306364 -0.95104897 0.31592367 0.43628453 -0.74905884 -0.17141096 -0.35913151 -0.11654172 -0.27070619 -0.05813532 -0.08783951]	-0.001388

Let's evaluate the performance of the best Lasso model that we were able to construct

```
In [91]: # Calculate the R squared for the training and validation datasets
lasso_r_squared_train = round(lasso_best.score(lasso_X_train, lasso_y_train), 3)
lasso_r_squared_validation = round(lasso_best.score(lasso_X_validation, lasso_y_validation), 3)
print(f"The R squared for the best Lasso model for the training set is: {lasso_r_squared_train}")
print(f"The R squared for the best Lasso model for the validation set is: {lasso_r_squared_validation}")

# Calculate the MSE of the Lasso model's predictions using the training and testing
lasso_y_predictions_train = lasso_model.predict(lasso_X_train)
lasso_mse_train = mean_squared_error(lasso_y_train, lasso_y_predictions_train)
print(f"The mean squared error of the Lasso Regression model's predictions using the training set is: {lasso_mse_train}")

lasso_y_predictions_validation = lasso_model.predict(lasso_X_validation)
lasso_mse_validation = mean_squared_error(lasso_y_validation, lasso_y_predictions_validation)
print(f"The mean squared error of the Lasso Regression model's predictions using the validation set is: {lasso_mse_validation}")
```

The R squared for the best Lasso model for the training set is: 0.851.

The R squared for the best Lasso model for the validation set is: 0.863.

The mean squared error of the Lasso Regression model's predictions using the log home sale prices in the training dataset is 144.98172527888673.

The mean squared error of the Lasso Regression model's predictions using the log home sale prices in the validation dataset is 146.1750136575638.

Rework LASSO analysis on full set of variables including encoded

```
In [92]: # Split the dataset into training and testing dataframes

scaler = StandardScaler()
x_raw = housing_training_data_large.select_dtypes(exclude=['object']).drop(columns = ['SalePrice'])
x_scale = xscaler.fit_transform(x_raw)
#y_scale = scaler.fit_transform(np.array(housing_training_data['SalePrice']).reshape(-1,))
y = housing_training_data_large['SalePrice']
new_lasso_X_train, new_lasso_X_validation, new_lasso_y_train, new_lasso_y_validation = \
    train_test_split(x_scale, y, test_size=0.2, random_state=42)

cv = KFold(n_splits=5, shuffle=True)

lasso_alphas = np.linspace(0, 0.02, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, x_scale, y_scale,
        cv, scoring = 'r2', model_name = 'LASSO',
        X_test = new_lasso_X_validation, y_test = new_lasso_y_validation,
        draw_plot = True, filename = 'lasso_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)
```


```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 35.31037983227483, tolerance: 0.1175439996562544
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 41.18845875616272, tolerance: 0.11994908768764215
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 37.21134079370813, tolerance: 0.11510572462008437
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 36.84929949601944, tolerance: 0.11130167392847584
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 37.827855076666005, tolerance: 0.11805824792012207
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:44: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 49.588607351090744, tolerance: 0.1455
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5323072608091621, tolerance: 0.1176629040603749
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5787914854352749, tolerance: 0.12076187790379247
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.3235088258803529, tolerance: 0.11008462468197588
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5329690477623359, tolerance: 0.11636219756433114
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.521069772902834, tolerance: 0.11705095348674897
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6918563681531253, tolerance: 0.1455
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.612771358713502, tolerance: 0.12034767772202624
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5947880036551112, tolerance: 0.1132149441195675
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.62415284785
```



```
41554, tolerance: 0.11364270760054555
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6327147653322243, tolerance: 0.11675327959970194
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6681823291409472, tolerance: 0.11799409528077164
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6514220902216863, tolerance: 0.1455
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8947401752448201, tolerance: 0.11705600642769155
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.7026529883925292, tolerance: 0.11494023970738529
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.720739335856166, tolerance: 0.12229740025743394
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.7235079052683488, tolerance: 0.10953361819285792
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.7497113939463134, tolerance: 0.11802449832063418
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.926633450807131, tolerance: 0.1455
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.2949261799653158, tolerance: 0.1127727010985301
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8613286977222856, tolerance: 0.12101298580921589
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.78439334450
```

```
01596, tolerance: 0.11781216750580778
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.7676279311636591, tolerance: 0.11565185578018948
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8710950709891137, tolerance: 0.11464783456618256
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.0437834903535048, tolerance: 0.1455
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8883606659324386, tolerance: 0.11931861656109605
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.9503544935621306, tolerance: 0.11931983307743459
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.913048650344507, tolerance: 0.11420908182302415
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5323298067461479, tolerance: 0.11525065812982799
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.510684336506273, tolerance: 0.1455
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.14887124087974257, tolerance: 0.11038849436093542
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.4381716007762151, tolerance: 0.10741983918266827
positive)
```

Chosen alpha: 0.00960

Validation score: 0.90266

Test score at chosen alpha: -5.52195

In [95]: `### BASED ON LASSOCV:`

```
# Show best value of tuning parameter (alpha) chosen by cross validation
print(f"The best value of the tuning parameter, alpha, chosen by cross-validation is {
# Set best tuning parameter (alpha) and use it to fit our final Lasso regression model
lasso_best = Lasso(alpha=lasso_model.alpha_)
lasso_best.fit(x_scale, y_scale)

lasso_best.coef_
```

The best value of the tuning parameter, alpha, chosen by cross-validation is 0.0035778259276540722.

Out[95]: `Lasso(alpha=0.0035778259276540722, copy_X=True, fit_intercept=True,
max_iter=1000, normalize=False, positive=False, precompute=False,
random_state=None, selection='cyclic', tol=0.0001, warm_start=False)`

```

Out[95]: array([-0.00000000e+00, -4.11641928e-02,  1.68674529e-02,  7.24623204e-02,
  1.42221206e-01,  6.37091222e-02,  6.65378779e-02,  6.44723957e-03,
  5.91435950e-02,  4.70831383e-02,  4.33865364e-02,  9.97823843e-02,
  0.00000000e+00, -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
  0.00000000e+00, -1.02289184e-02,  2.09573638e-01,  3.33189236e-03,
 -5.46213183e-03,  1.39061621e-03,  1.71937649e-02, -4.92538653e-02,
 -2.9555253e-02,  5.11046900e-02,  4.97189772e-02,  1.15072485e-02,
 -0.00000000e+00,  0.00000000e+00,  3.68962259e-02,  4.41013884e-02,
  1.83428461e-02, -0.00000000e+00,  2.18563071e-02,  1.96597065e-02,
  0.00000000e+00,  9.05923345e-03,  2.55688945e-02,  7.29072340e-03,
  0.00000000e+00, -9.49453220e-03,  0.00000000e+00,  1.21863168e-02,
  0.00000000e+00,  8.79282952e-04,  1.38182045e-02, -0.00000000e+00,
  1.78978199e-01, -1.88381909e-02,  5.29227611e-03,  0.00000000e+00,
  0.00000000e+00, -1.65701625e-02, -5.31651431e-03,  1.08271695e-02,
 -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  9.79819716e-03,
 -1.36340235e-02, -0.00000000e+00,  8.64029619e-03, -2.56381400e-17,
 -1.38794294e-03,  2.78942011e-02, -7.66727125e-03, -2.95661698e-03,
  0.00000000e+00, -0.00000000e+00,  3.91229824e-03, -2.73234175e-02,
  0.00000000e+00,  7.81750228e-04,  1.26543881e-02,  2.66377844e-02,
  0.00000000e+00, -8.56622265e-04,  4.76431492e-02, -3.94484436e-03,
 -7.05947109e-04,  0.00000000e+00, -0.00000000e+00, -5.91519048e-03,
 -3.32237844e-03,  1.44601176e-02, -9.20523700e-03,  4.60277418e-02,
  9.21870562e-02, -0.00000000e+00, -0.00000000e+00,  8.47176710e-03,
  0.00000000e+00,  2.18324747e-02,  7.29085570e-02, -0.00000000e+00,
  7.89579560e-03, -1.37850752e-03, -5.16928423e-04,  3.29263639e-02,
 -0.00000000e+00,  0.00000000e+00, -1.44121563e-02,  9.93173579e-04,
 -0.00000000e+00,  1.99300323e-04,  0.00000000e+00, -0.00000000e+00,
 -0.00000000e+00,  8.20307862e-03, -5.93864085e-03, -8.25542450e-03,
  0.00000000e+00,  4.44285365e-03,  4.81227342e-02,  1.43636382e-02,
 -0.00000000e+00, -1.11366489e-02, -0.00000000e+00, -5.42761342e-03,
  2.08204330e-03,  0.00000000e+00, -3.26274640e-03, -2.84325953e-03,
 -0.00000000e+00, -3.34462380e-04,  9.68455305e-04, -0.00000000e+00,
 -1.53003023e-02, -0.00000000e+00,  1.50152661e-03,  0.00000000e+00,
  3.06431925e-03,  0.00000000e+00,  1.39802794e-02,  4.52769348e-03,
 -1.02443687e-03, -9.93173127e-03, -1.17115017e-03,  2.19971605e-02,
  0.00000000e+00, -1.29534324e-03, -0.00000000e+00,  2.47370263e-02,
 -0.00000000e+00,  4.66886522e-03, -1.06781329e-02, -0.00000000e+00,
  2.44173254e-03, -6.99740581e-03, -0.00000000e+00,  0.00000000e+00,
  0.00000000e+00, -2.17791584e-02,  0.00000000e+00,  0.00000000e+00,
 -0.00000000e+00,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
  1.36015094e-02, -5.47260759e-03, -0.00000000e+00,  0.00000000e+00,
 -9.48333527e-03, -9.08224695e-03, -1.88499913e-03, -0.00000000e+00,
  0.00000000e+00,  7.27601399e-03, -3.62226615e-03, -1.13712099e-02,
 -2.38261473e-02,  0.00000000e+00,  1.06676924e-02,  1.07326229e-03,
  4.78343573e-03, -3.61537054e-03,  0.00000000e+00,  0.00000000e+00,
 -8.16079028e-03, -0.00000000e+00,  0.00000000e+00,  0.00000000e+00,
  0.00000000e+00, -1.15543355e-02, -1.99967616e-03, -0.00000000e+00,
  2.98255402e-02,  1.36725674e-02,  0.00000000e+00,  0.00000000e+00,
  6.97761132e-02, -1.34062669e-03, -2.05711341e-02,  0.00000000e+00,
 -0.00000000e+00,  0.00000000e+00,  1.61264655e-02, -5.88018109e-03,
  1.62512225e-03, -9.02280592e-03,  0.00000000e+00,  1.70625401e-02,
 -0.00000000e+00,  1.20341614e-02, -1.80073545e-05,  0.00000000e+00,
  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,  0.00000000e+00,
 -0.00000000e+00,  0.00000000e+00, -1.47006792e-02,  2.41276733e-03,
  0.00000000e+00,  0.00000000e+00,  1.38944991e-03, -7.61111234e-04,
  1.79173217e-03, -3.33226841e-03, -8.72097534e-03, -6.71160593e-03,
 -0.00000000e+00,  0.00000000e+00, -3.79393444e-03, -1.37204495e-02,
  3.87021485e-02, -1.07768334e-02, -8.66479352e-03, -0.00000000e+00,
  1.73038571e-02,  0.00000000e+00,  0.00000000e+00,  3.07172169e-02,
  0.00000000e+00,  3.11442357e-15, -2.69444207e-02, -0.00000000e+00,

```

```
-5.11827731e-03,  5.05126470e-03,  7.91052204e-03,  1.78206443e-03,
 0.00000000e+00,  0.00000000e+00,  6.37785659e-02,  2.83827499e-04,
-8.16526918e-03, -2.33267870e-02,  0.00000000e+00,  0.00000000e+00,
-9.45648888e-03,  0.00000000e+00,  0.00000000e+00])
```

REPEAT WITH RIDGE

```
In [96]: # Split the dataset into training and testing dataframes


scaler = StandardScaler()
x_raw = housing_training_data.select_dtypes(exclude=['object']).drop(columns = ['SalePrice'])
x_scale = xscale.fit_transform(x_raw)
y_scale = np.array(housing_training_data['SalePrice']).reshape(-1,1)

new_ridge_X_train, new_ridge_X_validation, new_ridge_y_train, new_ridge_y_validation = \
                                                    y_scale, train_indices, test_indices, random_state

cv = KFold(n_splits=5, shuffle=True)

lasso_alphas = np.linspace(0, 0.02, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, new_ridge_X_train, new_ridge_y_train,
        cv, scoring = 'r2', model_name = 'Ridge',
        X_test = new_ridge_X_validation, y_test = new_ridge_y_validation,
        draw_plot = True, filename = 'ridge_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)
```



Chosen alpha: 0.01900

Validation score: 0.86132

Test score at chosen alpha: 0.88396

REPEAT WITH RIDGE USING SUBNET OF ENCODED VARIABLE SET

```
In [97]: scaler = StandardScaler()
x_raw = housing_training_data_large_subset.select_dtypes(exclude=['object']).drop(columns=['SalePrice'])
x_scale = xscale.fit_transform(x_raw)
y_scale = np.array(housing_training_data['SalePrice']).reshape(-1,1)

new_ridge_X_train, new_ridge_X_validation, new_ridge_y_train, new_ridge_y_validation = \
    ridge(X_train=x_scale, X_validation=x_scale, y_train=y_scale, y_validation=y_scale,
          random_state=42)

cv = KFold(n_splits=5, shuffle=True)


lasso_alphas = np.linspace(0, 1, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, new_ridge_X_train, new_ridge_y_train,
        cv, scoring = 'r2', model_name = 'Ridge',
```

```

X_test = new_ridge_X_validation, y_test = new_ridge_y_validation,
draw_plot = True, filename = 'ridge_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)

```



Chosen alpha: 0.26000
 Validation score: 0.87303
 Test score at chosen alpha: 0.88549

```

In [118... # Show best Ridge model coefficients and names
regr = Ridge(alpha=chosen_alpha)
r_model = regr.fit(x_scale, y_scale)

r_model.coef_

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-118-d679e1c0ac7f> in <module>
      1 # Show best Ridge model coefficients and names
      2 regr = Ridge(alpha=chosen_alpha)
----> 3 r_model = regr.fit(x_scale, y_scale)
      4
      5 r_model.coef_

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_ridge.py in fit(self, X, y, sample_weight)
    764         self : returns an instance of self.
    765         """
--> 766         return super().fit(X, y, sample_weight=sample_weight)
    767
    768

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_ridge.py in fit(self, X, y, sample_weight)
    545         accept_sparse=_accept_sparse,
    546         dtype=_dtype,
--> 547         multi_output=True, y_numeric=True)
    548         if sparse.issparse(X) and self.fit_intercept:
    549             if self.solver not in ['auto', 'sparse_cg', 'sag']:

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/utils/validation.py in check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ensure_min_features, y_numeric, warn_on_dtype, estimator)
    763         y = y.astype(np.float64)
    764
--> 765         check_consistent_length(X, y)
    766
    767         return X, y

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/utils/validation.py in check_consistent_length(*arrays)
    210         if len(uniques) > 1:
    211             raise ValueError("Found input variables with inconsistent numbers of"
--> 212                             " samples: %r" % [int(l) for l in lengths])
    213
    214

ValueError: Found input variables with inconsistent numbers of samples: [1459, 1455]

```

Try ElasticNet...

```

In [99]: # Split the dataset into training and testing dataframes

elastic_scaler = StandardScaler()
x_raw = housing_training_data.select_dtypes(exclude=['object']).drop(columns = ['SalePrice'])
x_scale = elastic_scaler.fit_transform(x_raw)
y_scale = elastic_scaler.fit_transform(np.array(housing_training_data['SalePrice']).res

new_elastic_X_train, new_elastic_X_validation, new_elastic_y_train, new_elastic_y_vali
                                                    y_scale, t
                                                    random_sta

cv = KFold(n_splits=5, shuffle=True)

```

```
lasso_alphas = np.linspace(0, 0.02, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, new_elastic_X_train, new_elastic_y_train,
        cv, scoring = 'r2', model_name = 'ElasticNet',
        X_test = new_elastic_X_validation, y_test = new_elastic_y_validation,
        draw_plot = True, filename = 'EN_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 49.312537538118214, tolerance: 0.08376794942129816
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 51.480662924915535, tolerance: 0.08296929794930137
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 50.30861021939106, tolerance: 0.0832879990067399
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 54.22892836581726, tolerance: 0.08423980593987467
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
```



```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 47.956581514630635, tolerance: 0.08064710497810411
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:44: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 64.21556748262068, tolerance: 0.10374553746774125
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 14.832410968152516, tolerance: 0.07984098495757876
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 17.271489777890878, tolerance: 0.08478563025384582
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 18.29386176899299, tolerance: 0.08573864849881115
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 17.145604377213814, tolerance: 0.08156778779801227
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 10.596691086372644, tolerance: 0.08301386036646154
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 21.24874758046724, tolerance: 0.10374553746774125
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5892763164067958, tolerance: 0.08031832214608478
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.336113848486363, tolerance: 0.087668087863972
    positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.72855298349
```


```
26194, tolerance: 0.08761967961525664
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.591400171377643, tolerance: 0.08012040170117148
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.094574612040148, tolerance: 0.07922056992737007
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 2.121453235503438, tolerance: 0.10374553746774125
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.5014416931059742, tolerance: 0.08387924134623305
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.4545618782720169, tolerance: 0.08478729028699415
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8605940139759554, tolerance: 0.0829635134454385
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.3115864298690951, tolerance: 0.07829629426425311
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.1566717574116296, tolerance: 0.08500518242117788
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.919223964472991, tolerance: 0.10374553746774125
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.2198442536516296, tolerance: 0.07805887084269865
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6333163430416988, tolerance: 0.07836715183096404
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.72355389023
```

```
48395, tolerance: 0.08657749475035348
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.062481431203679, tolerance: 0.08943469232554455
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.6492670328387078, tolerance: 0.10374553746774125
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5461789596847098, tolerance: 0.08285161879543186
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6481802451306606, tolerance: 0.08645037608406281
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.1663117112920673, tolerance: 0.08853233855876791
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8666358466268065, tolerance: 0.07659656457777933
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6620399149314551, tolerance: 0.0804650902509035
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.2426757080564101, tolerance: 0.10374553746774125
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.8474800493402341, tolerance: 0.08409565251635383
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.7024958693655918, tolerance: 0.08326943437755133
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.6901345303540012, tolerance: 0.08759098461862767
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.92795236143
```

```

25171, tolerance: 0.10374553746774125
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5999416446549617, tolerance: 0.08124478175214664
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5289558535225183, tolerance: 0.08421883793675204
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.18427547147619805, tolerance: 0.08514750485032266
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.14913533380695299, tolerance: 0.08255240780472282
positive)

```



Chosen alpha: 0.01280

Validation score: 0.86336

Test score at chosen alpha: 0.88575

REPEAT WITH ELASTIC NET USING FULL ENCODED VARIABLE SET with LOG TRANSFORMED SALES PRICE

```
In [100... # Split the dataset into training and testing dataframes

elasticscaler = StandardScaler()
x_raw = housing_training_data_large_common.select_dtypes(exclude=['object'])
x_scale = xscaler.fit_transform(x_raw)
y_scale = elasticscaler.fit_transform(np.array(np.log(housing_training_data['SalePrice'])))

new_elastic_X_train, new_elastic_X_validation, new_elastic_y_train, new_elastic_y_validation = \
    train_test_split(x_scale, y_scale, test_size=0.2, random_state=42)

cv = KFold(n_splits=5, shuffle=True)

lasso_alphas = np.linspace(0, 0.02, 101)

chosen_alpha, max_validation_score, test_score_at_chosen_alpha = \
    regmodel_param_test(
        lasso_alphas, new_elastic_X_train, new_elastic_y_train,
        cv, scoring = 'r2', model_name = 'ElasticNet',
        X_test = new_elastic_X_validation, y_test = new_elastic_y_validation,
        draw_plot = True, filename = 'EN_wide_search')
print("Chosen alpha: %.5f" % \
      chosen_alpha)
print("Validation score: %.5f" % \
      max_validation_score)
print("Test score at chosen alpha: %.5f" % \
      test_score_at_chosen_alpha)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 16.54411307026035, tolerance: 0.07539341997255294
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 16.869904052605243, tolerance: 0.07725168572866303
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 16.70321159594418, tolerance: 0.0788264889258652
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 15.094661782926801, tolerance: 0.07603936850041647
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/model_selection/_validation.py:515: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
estimator.fit(X_train, y_train, **fit_params)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
```




```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 17.191931226457065, tolerance: 0.07694340945238295
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/ipykernel_launcher.py:44: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 22.477966273266922, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.9039995259503897, tolerance: 0.07623876165913449
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4.0645798620300795, tolerance: 0.07610252455058931
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 6.4344266030482355, tolerance: 0.0759868798099048
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 8.315870492235138, tolerance: 0.08114298606601994
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 1.4756700923142212, tolerance: 0.0749816143695759
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 3.1191442318222755, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.17001955667550916, tolerance: 0.07601276646072924
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.9584273161508392, tolerance: 0.0777792267773953
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.17253331035
```

```
206187, tolerance: 0.07360550635910844
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.4690948303690341, tolerance: 0.07955760019098691
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5694623396295739, tolerance: 0.07742283223216413
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.2251241112543383, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.21804421017856157, tolerance: 0.07877312557424299
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.4317609640254183, tolerance: 0.07947117074664821
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.5706252372928873, tolerance: 0.0758603512428788
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.22306834893089444, tolerance: 0.07425728258752042
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.87029802379476, tolerance: 0.07612537381325887
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.22437728195586004, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.16971187968005808, tolerance: 0.07634739405347273
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.14486588347944718, tolerance: 0.07332520826459797
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.54815887425
```



```
55985, tolerance: 0.08151762677599025
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.16442160677052797, tolerance: 0.07554830737803726
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.34365312008965887, tolerance: 0.07776746432901534
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.21164941245021396, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.14959850747870718, tolerance: 0.07278342135686353
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.11671506983923763, tolerance: 0.07719096820309204
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.1528535392859922, tolerance: 0.08063668034245289
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.34636192639790764, tolerance: 0.07820399985909574
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.18532742750152664, tolerance: 0.09613331437474897
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.11294847871543112, tolerance: 0.07960459611255288
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.3371327535362276, tolerance: 0.07744520486830796
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.10990394275642146, tolerance: 0.0796201418006938
positive)
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:476: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.14075228147
```

40787, tolerance: 0.07584299507683794
positive)



Chosen alpha: 0.01620

Validation score: 0.92317

Test score at chosen alpha: 0.90052

```
In [101... # Let's elastic build model with this alpha
regr = ElasticNet(alpha=chosen_alpha)
elastic_model = regr.fit(x_scale, y_scale)

ElasticNet(random_state=1)
print(regr.coef_)
print(regr.intercept_)

y_pred = elastic_model.predict(x_scale)

plt.subplot(1, 2, 1)
plt.scatter(y_scale, y_pred)
plt.xlabel('Sale Price (Standard Scaled)')
plt.ylabel('Predicted Sale Price')
plt.title('ElasticNet: Sale Price vs Predicted Sale Price')

# Residuals

plt.subplot(1, 2, 2)
```

```
sns.residplot(x=y_scale, y=y_pred)
plt.xlabel('Sale Price')
plt.ylabel('Residual')
plt.title('Residual Plot')

MSE = np.square(np.subtract(y_scale,y_pred)).mean()
print("MSE:",MSE)

r2 = r2_score(y_scale, y_pred)
print("R_sq:",r2)
```

```
Out[101]: ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,
                    max_iter=1000, normalize=False, positive=False, precompute=False,
                    random_state=1, selection='cyclic', tol=0.0001, warm_start=False)
```

[-1.25799231e-03	-0.00000000e+00	2.52828503e-02	3.99026969e-02
1.76904467e-01	9.13294632e-02	1.02098694e-01	0.00000000e+00
2.96895048e-03	1.61808229e-02	2.42243790e-02	6.95349451e-02
0.00000000e+00	-0.00000000e+00	0.00000000e+00	2.51643279e-03
0.00000000e+00	-1.86622997e-05	1.84006566e-01	2.46860056e-02
0.00000000e+00	1.01755101e-02	1.72673876e-02	-0.00000000e+00
-1.65936356e-02	3.32505813e-02	4.93259348e-03	2.61716317e-02
1.52140017e-02	0.00000000e+00	5.21161155e-02	4.10254337e-02
2.70558643e-02	0.00000000e+00	2.30652808e-02	1.34023067e-02
1.67813390e-03	1.34991981e-03	2.39006708e-02	0.00000000e+00
-0.00000000e+00	-0.00000000e+00	-0.00000000e+00	9.41176671e-04
3.21902052e-02	6.91241163e-03	2.64275356e-02	2.00981720e-01
-4.01255451e-02	-7.54500592e-02	3.79073812e-03	0.00000000e+00
0.00000000e+00	-5.01522928e-02	0.00000000e+00	3.22639563e-03
-0.00000000e+00	-0.00000000e+00	-0.00000000e+00	0.00000000e+00
-0.00000000e+00	-0.00000000e+00	0.00000000e+00	0.00000000e+00
1.84472801e-02	-0.00000000e+00	-0.00000000e+00	-0.00000000e+00
-0.00000000e+00	0.00000000e+00	-3.97772115e-03	-0.00000000e+00
-0.00000000e+00	-4.91541755e-04	1.13170266e-02	8.77678553e-03
-0.00000000e+00	4.85363737e-02	-1.17670065e-02	0.00000000e+00
-0.00000000e+00	-2.15960818e-02	-1.17832513e-03	-0.00000000e+00
0.00000000e+00	-0.00000000e+00	3.80998842e-03	2.74290398e-02
-2.23187875e-03	-0.00000000e+00	-0.00000000e+00	-0.00000000e+00
1.38466880e-02	2.84149704e-02	0.00000000e+00	3.71488597e-03
-1.27563708e-02	-0.00000000e+00	2.59726671e-02	-0.00000000e+00
0.00000000e+00	-9.13520856e-03	0.00000000e+00	-0.00000000e+00
0.00000000e+00	-0.00000000e+00	0.00000000e+00	0.00000000e+00
0.00000000e+00	-0.00000000e+00	2.51263310e-02	0.00000000e+00
-6.15091942e-04	-1.55519981e-02	-0.00000000e+00	0.00000000e+00
0.00000000e+00	-0.00000000e+00	0.00000000e+00	0.00000000e+00
-0.00000000e+00	6.55827650e-04	0.00000000e+00	-1.29595794e-04
-0.00000000e+00	0.00000000e+00	0.00000000e+00	0.00000000e+00
-0.00000000e+00	-0.00000000e+00	-0.00000000e+00	7.28270473e-04
-0.00000000e+00	0.00000000e+00	-1.33082845e-02	2.41896524e-02
-0.00000000e+00	0.00000000e+00	-6.11469164e-03	5.43587585e-03
-0.00000000e+00	0.00000000e+00	0.00000000e+00	-1.15170483e-02
-0.00000000e+00	-0.00000000e+00	0.00000000e+00	-0.00000000e+00
-0.00000000e+00	-0.00000000e+00	0.00000000e+00	-0.00000000e+00
-0.00000000e+00	0.00000000e+00	-0.00000000e+00	-0.00000000e+00
-0.00000000e+00	0.00000000e+00	0.00000000e+00	-0.00000000e+00
-2.85742997e-03	0.00000000e+00	-0.00000000e+00	1.11322540e-02
4.43071426e-04	-0.00000000e+00	-0.00000000e+00	-0.00000000e+00
6.20371380e-03	-0.00000000e+00	-0.00000000e+00	2.62451044e-02
0.00000000e+00	6.22903279e-03	-4.77455788e-03	-7.79837630e-03
0.00000000e+00	0.00000000e+00	0.00000000e+00	0.00000000e+00
0.00000000e+00	3.04113582e-02	0.00000000e+00	-1.07545621e-02
0.00000000e+00	0.00000000e+00	0.00000000e+00	8.97144497e-03
-0.00000000e+00	0.00000000e+00	-0.00000000e+00	-2.46339358e-04
3.10139060e-03	-3.50975469e-03	4.64976308e-03	-0.00000000e+00
0.00000000e+00	0.00000000e+00	-0.00000000e+00	-0.00000000e+00
7.71655985e-03	-1.56874171e-02	0.00000000e+00	0.00000000e+00
0.00000000e+00	-0.00000000e+00	-0.00000000e+00	-7.55102047e-03
-2.44924883e-02	0.00000000e+00	0.00000000e+00	-9.78471378e-03
-1.30260125e-02	2.53278389e-02	-5.09277918e-03	3.68371423e-03
-0.00000000e+00	0.00000000e+00	-0.00000000e+00	-0.00000000e+00
-0.00000000e+00	3.26300960e-03	-0.00000000e+00	0.00000000e+00
-6.50114630e-04	-0.00000000e+00	0.00000000e+00	8.86416409e-04
6.24797758e-03	-0.00000000e+00	0.00000000e+00	2.68387209e-02
0.00000000e+00	-8.86138386e-03	-4.00799608e-02	0.00000000e+00

```
-0.00000000e+00 -1.30035001e-02 0.00000000e+00 0.00000000e+00]
[-2.18349378e-15]
```

```
Out[101]: <AxesSubplot:>
```

```
Out[101]: <matplotlib.collections.PathCollection at 0x7fee8a0ef850>
```

```
Out[101]: Text(0.5, 0, 'Sale Price (Standard Scaled)')
```

```
Out[101]: Text(0, 0.5, 'Predicted Sale Price')
```

```
Out[101]: Text(0.5, 1.0, 'ElasticNet: Sale Price vs Predicted Sale Price')
```

```
Out[101]: <AxesSubplot:>
```

```
Out[101]: <AxesSubplot:>
```


```
Out[101]: Text(0.5, 0, 'Sale Price')
```

```
Out[101]: Text(0, 0.5, 'Residual')
```

```
Out[101]: Text(0.5, 1.0, 'Residual Plot')
```

```
MSE: 1.8975655358794232
```

```
R_sq: 0.93603420818627
```




Let's check the coefficients to determine the important predictors in the model

```
In [102... coef = pd.Series(elastic_model.coef_, index = x_raw.columns)
important_features = pd.concat([coef.sort_values().head(10),
                                coef.sort_values().tail(10)])
important_features.plot(kind = "barh")
plt.title("Coefficients in the ElasticNet Model")
```

```
Out[102]: <AxesSubplot:>
```

```
Out[102]: Text(0.5, 1.0, 'Coefficients in the ElasticNet Model')
```



TotalSF, GrLivArea, and Overall Quality appear to have strong, positive effects on Sales Price.

Try using LassoCV

```
In [103]: from sklearn.linear_model import LassoCV
from sklearn.datasets import make_regression
reg = LassoCV(cv=5, random_state=0).fit(x_scale, y_scale)
reg.score(x_scale, y_scale)
y_pred = reg.predict(x_scale)
print(reg.alpha_)
```

```
/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:1088: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
Out[103]: 0.9346873914793059
0.009417743973449105
```

Running Models on Test Data for Kaggle Predictions

Scale x-values and transform using PCA, then use those values in the ridge model we created above to predict values of Sales Price

```
In [104]: # select numeric variables in the test data for x
x_test_raw = housing_testing_data.select_dtypes(exclude=['object'])
# scale x-values
x_test_scale = xscaler.fit_transform(x_test_raw)
# transform x-values using PCA
x_test_pca_scale = pca.fit_transform(x_test_scale)
# use x-values to predict y values
scaled_y_predtest = ridgemodel.predict(x_test_pca_scale)

# Apply inverse scaling transformation to predicted y-values
y_predtest = yscaler.inverse_transform(scaled_y_predtest)
# Exponential transform the predictions (since the model output log transformed y-values)
y_pred_final = np.exp(y_predtest)
```

Create dataframe with Id and Sales Price predictions

```
In [105... # Create a dataframe with the y predictions
predictiondf=pd.DataFrame(y_pred_final, columns=['SalePrice'])
# Add the Id column to the front of the dataframe
predictiondf.insert(0, 'Id', housing_testing_data['Id'])

predictiondf.head()
```

```
Out[105]:
```

	Id	SalePrice
0	1461	256784.333872
1	1462	203132.211028
2	1463	161976.051358
3	1464	145088.338480
4	1465	148360.496704

```
In [106... #output predictions to csv
predictiondf.to_csv('test_salespriceridge_v6_ridge.csv', index=False)
```

Kaggle Results - Ridge Regression

Upon submission of our home price predictions from the ridge regression model into Kaggle using Claire Markey's username, the team achieved a RMSE (as calculated using the log of the predicted and actual home prices) of 0.16954 for the testing dataset as displayed in the screenshot from the Kaggle leaderboard below.

```
In [107... plt.figure(figsize = (15, 15))
kaggle_results = plt.imread('Kaggle_RMSE_ridge.jpg')
plt.imshow(kaggle_results)
plt.axis("off")
plt.show()
```

```
Out[107]: <Figure size 1080x1080 with 0 Axes>
```

```
Out[107]: <matplotlib.image.AxesImage at 0x7fee7045d650>
```

```
Out[107]: (-0.5, 1246.5, 126.5, -0.5)
```



test_salespriceridge.csv
Complete · 2m ago · ridge v2

0.16954

Predicting Home Sale Prices In the Testing Dataset using the Polynomial Model

Let's predict Sales Price using our polynomial model

```
In [108... poly_model = PolynomialFeatures(degree=degree)

# transform out polynomial features
poly_x_test_values = poly_model.fit_transform(x_test_pca_scale)

y_poly_pred = poly_regression_model.predict(poly_x_test_values)

# Apply inverse scaling transformation to predicted y-values
y_poly_predtest = ypolyscaler.inverse_transform(y_poly_pred)
```

```
In [109... # Create a dataframe with the y predictions
predictionpolydf=pd.DataFrame(y_poly_predtest, columns=['SalePrice'])
# Add the Id column to the front of the dataframe
predictionpolydf.insert(0, 'Id', housing_testing_data['Id'])
predictionpolydf.head()
```

```
Out[109]:
```

	Id	SalePrice
0	1461	261902.092805
1	1462	205879.875005
2	1463	155206.325441
3	1464	135215.922537
4	1465	141244.457630

```
In [110... #output predictions to csv

predictionpolydf.to_csv('test_salespricepoly_v5.csv', index=False)
```

Kaggle Results - Polynomial Regression

Upon submission of our home price predictions from the polynomial model into Kaggle using Claire Markey's username, the team achieved a RMSE (as calculated using the log of the predicted and actual home prices) of 0.163 for the testing dataset as displayed in the screenshot from the Kaggle leaderboard below.

```
In [111... plt.figure(figsize = (15, 15))
kaggle_results = plt.imread('Kaggle_RMSE_poly.jpg')
plt.imshow(kaggle_results)
plt.axis("off")
plt.show()
```

```
Out[111]: <Figure size 1080x1080 with 0 Axes>
```


```
Out[111]: <matplotlib.image.AxesImage at 0x7fee6f16a2d0>
```

```
Out[111]: (-0.5, 1300.5, 420.5, -0.5)
```


Overview	Data	Code	Discussion	Leaderboard	Rules	Team	Submissions	Submit Predictions	...
----------	------	------	------------	-------------	-------	------	-------------	--------------------	-----

Submissions

All
Successful
Errors
Recent ▾

Submission and Description	Public Score ⓘ
 test_salespricepoly_v5.csv Complete · now · poly v3	0.163

Kaggle Results - Lasso Regression

```
In [112...] # Set up Test dataframe for Lasso Regression Analysis

lasso_X_test = housing_testing_data.copy(deep=True)

lasso_X_test = lasso_X_test[ ['LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
                             'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
                             'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'KitchenArea',
                             'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'ScreenPorch',
                             'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']]

from sklearn.preprocessing import StandardScaler

lasso_X_test[numerical_predictors] = lasso_scaler.transform(lasso_X_test[numerical_predictors])

In [113...] # Exponential transform the predictions (since the model output log transformed y-values)
lasso_y_predictions_test = lasso_model.predict(lasso_X_test)
y_predictions_final_lasso = np.exp(lasso_y_predictions_test)

# Create a dataframe with the y predictions
predictiondf_lasso=pd.DataFrame(y_predictions_final_lasso, columns=['SalePrice'])
# Add the Id column to the front of the dataframe
predictiondf_lasso.insert(0, 'Id', housing_testing_data['Id'])

#output predictions to csv
predictiondf_lasso.to_csv('test_sales_price_lasso_v1.csv', index=False)


In [114...] # Display the kaggle results associated with the Lasso Regression Model
plt.figure(figsize = (15, 15))
kaggle_results = plt.imread('Kaggle_RMSE_lasso.jpg')
plt.imshow(kaggle_results)
plt.axis("off")
plt.show()
```

Out[114]: <Figure size 1080x1080 with 0 Axes>

Out[114]: <matplotlib.image.AxesImage at 0x7fee6f4b4390>

Out[114]: (-0.5, 1478.5, 335.5, -0.5)

Submissions

Submission and Description		Public Score ⓘ
<div>  test_sales_price_lasso_v1.csv Complete · 1m ago · Predictions from Lasso Regression Model </div>		0.16194

Kaggle Results - Elastic Regression

```
In [115... x_raw = housing_testing_data_large_common.select_dtypes(exclude=['object'])
x_scale = xscaler.fit_transform(x_raw)

# predict Sales Price
y_scale_elastic_pred = elastic_model.predict(x_scale)

# transform the predictions (since the model output scaled y-values)
y_log_elastic_pred = elasticscaler.inverse_transform(y_scale_elastic_pred)
y_elastic_pred = np.exp(y_log_elastic_pred)
y_elastic_pred
```

```
Out[115]: array([124133.94005013, 160183.6995892 , 179493.67266664, ...,
        172452.63034495, 120158.64228585, 226165.83170092])
```

```
In [116... # Create a dataframe with the y predictions
predictiondf_elastic=pd.DataFrame(y_elastic_pred, columns=['SalePrice'])
# Add the Id column to the front of the dataframe
predictiondf_elastic.insert(0, 'Id', housing_testing_data['Id'])

#output predictions to csv
predictiondf_elastic.to_csv('test_sales_price_elastic_v7.csv', index=False)
```

```
In [117... # Display the kaggle results associated with the Lasso Regression Model
plt.figure(figsize = (15, 15))
kaggle_results = plt.imread('Kaggle_RMSE_elastic_full.jpg')
plt.imshow(kaggle_results)
plt.axis("off")
plt.show()
```

```
Out[117]: <Figure size 1080x1080 with 0 Axes>
```

```
Out[117]: <matplotlib.image.AxesImage at 0x7fee8b1a4390>
```

```
Out[117]: (-0.5, 1187.5, 121.5, -0.5)
```

<div>  test_sales_price_elastic_v7.csv Complete · now </div>		0.1356
---	--	---------------

```
In [ ]:
```